

**Batch: E2              Roll No.: 16010123325**

**Experiment No. 6**

**Title: Data Cleaning, Transformation and Feature Engineering Techniques**

Aim : Handle missing values, detect outliers, perform one-hot encoding, label encoding, feature scaling (min-max and z-score normalization), and generate polynomial features.

**Course Outcome:**

**CO3 :** Learn data cleaning, transformation, and feature engineering techniques.

**Books/ Journals/ Websites referred:**

1. [The Comprehensive R Archive Network](#)
2. [Posit](#)

**Resources used:**

**(Students should write the datasets used)**

---

**Theory:**

**Implementation:**

**A) Handling Missing Values**

In this example we will use the open repository of plants classification Iris.

**1) Data Loading**

```
> data()
> str(iris_data)
Error: object 'iris_data' not found
> str(iris)
'data.frame': 150 obs. of 5 variables:
 $ Sepal.Length: num 5.1 4.9 4.7 4.6 5 5.4 4.6 5 4.4 4.9 ...
 $ Sepal.Width : num 3.5 3 3.2 3.1 3.6 3.9 3.4 3.4 2.9 3.1 ...
 $ Petal.Length: num 1.4 1.4 1.3 1.5 1.4 1.7 1.4 1.5 1.4 1.5 ...
 $ Petal.Width : num 0.2 0.2 0.2 0.2 0.2 0.4 0.3 0.2 0.2 0.1 ...
 $ Species     : Factor w/ 3 levels "setosa","versicolor",...: 1 1 1 1 1 1 1 1 1 1 ...
```

## 2) Introducing missing values in the dataset

```
> iris_copy <- iris
>
> iris_copy$Sepal.Length[c(15, 20, 50, 67, 97, 118)] <- NA
> iris_copy$Sepal.Width[c(4, 80, 97, 106)] <- NA
> iris_copy$Petal.Length[c(5, 17, 35, 49)] <- NA
> summary(iris_copy)
   Sepal.Length   Sepal.Width   Petal.Length   Petal.Width      Species
Min.    :4.300   Min.    :2.000   Min.    :1.000   Min.    :0.100   setosa    :5
1st Qu.:5.100   1st Qu.:2.800   1st Qu.:1.600   1st Qu.:0.300   versicolor:5
Median :5.800   Median :3.000   Median :4.400   Median :1.300   virginica :5
Mean   :5.844   Mean   :3.062   Mean   :1.778   Mean   :0.851
3rd Qu.:6.400   3rd Qu.:3.375   3rd Qu.:5.100   3rd Qu.:1.800
Max.   :7.900   Max.   :4.400   Max.   :6.900   Max.   :2.500
NA's    :6       NA's    :4       NA's    :4
```

## 3) Find NA's in dataset

# The first thing we can do is to ask if there is any missing value in our table

```
> length(which(is.na(iris_copy)))
[1] 14
```

# We can check that we introduced 14 missing values in the table

# There are several ways to identify rows containing NA's.

# First we will use the complete.cases function (check ?complete.cases for information)

# This function returns only rows without NA's. Putting ! in front of it we get only rows with NA's

```
> iris_NA <- iris_copy[!complete.cases(iris_copy), ]
> iris_NA
  Sepal.Length Sepal.Width Petal.Length Petal.Width Species
4           4.6        NA         1.5       0.2   setosa
5           5.0        3.6         NA       0.2   setosa
15          NA         4.0         1.2       0.2   setosa
17          5.4        3.9         NA       0.4   setosa
20          NA         3.8         1.5       0.3   setosa
35          4.9        3.1         NA       0.2   setosa
49          5.3        3.7         NA       0.2   setosa
50          NA         3.3         1.4       0.2   setosa
```

# We see that we have 13 rows with missing values on it

# Another way is to search for TRUE values in the is.na function

```
> iris_NA <- iris_copy[rowSums(is.na(iris_copy)) > 0, ]
> iris_NA
  Sepal.Length Sepal.Width Petal.Length Petal.Width Species
4           4.6        NA         1.5       0.2   setosa
5           5.0        3.6         NA       0.2   setosa
15          NA         4.0         1.2       0.2   setosa
17          5.4        3.9         NA       0.4   setosa
20          NA         3.8         1.5       0.3   setosa
```

# We obtain the same result. However the function complete cases is

#### 4) Removing rows containing NA's

# Depending on the case, there are different things we can do with NA's. However there is not a unique and general solution to this issue. If the missing value can be calculated directly using other columns the problem is solved.

# The first choice can be to just remove the rows containing NA's

```
> complete.cases(iris_copy)
 [1]  TRUE  TRUE  TRUE FALSE FALSE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE
[14] TRUE FALSE  TRUE FALSE  TRUE  TRUE FALSE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE
[27] TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE FALSE  TRUE  TRUE  TRUE  TRUE
[40] TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE FALSE FALSE  TRUE  TRUE
[53] TRUE  TRUE
[66] TRUE FALSE  TRUE  TRUE
[79] TRUE FALSE  TRUE  TRUE
[92] TRUE  TRUE  TRUE  TRUE  TRUE FALSE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE
[105] TRUE FALSE  TRUE  TRUE
[118] FALSE  TRUE  TRUE
[131] TRUE  TRUE
[144] TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE
```

```
> iris_clean <- iris_copy[complete.cases(iris_copy), ]
> length(which(is.na(iris_clean)))
[1] 0
```

# The number of missing values in this table is 0

## 5) Substituting NA's with numerical values

# In other cases we don't want to lose the information that we have in one row with missing values

# In this case we will substitute the missing value with a numerical value

# The first thing we can do is to introduce the mean of a column in a missing value

# However it's more safe to use the median because it's not affected by outliers

# However we should be careful as in this case it's more correct to introduce the mean for the proper species

# We should do it column by column

```
> iris_copy[is.na(iris_copy$Sepal.Length) & (iris_copy$Species == "setosa"), "Sepal.Length"] <- median(iris_copy$Sepal.Length[which(iris_copy$Species == "setosa")], na.rm = TRUE)
> iris_NA <- iris_copy[!complete.cases(iris_copy), ]
> iris_NA
   Sepal.Length Sepal.Width Petal.Length Petal.Width Species
4          4.6        NA         1.5       0.2    setosa
5          5.0        3.6        NA        0.2    setosa
17         5.4        3.9        NA        0.4    setosa
35         4.9        3.1        NA        0.2    setosa
49         5.3        3.7        NA        0.2    setosa
67         NA         3.0        4.5       1.5 versicolor
```

# Now we have removed 3 NA's. Only 11 left

```
> iris_copy[is.na(iris_copy$Sepal.Length) & (iris_copy$Species == "versicolor"), "Sepal.Length"] <- median(iris_copy$Sepal.Length[which(iris_copy$Species == "versicolor")], na.rm = TRUE)
> iris_NA <- iris_copy[!complete.cases(iris_copy), ]
> iris_NA
   Sepal.Length Sepal.Width Petal.Length Petal.Width Species
4          4.60        NA         1.5       0.2    setosa
5          5.00        3.6        NA        0.2    setosa
17         5.40        3.9        NA        0.4    setosa
35         4.90        3.1        NA        0.2    setosa
49         5.30        3.7        NA        0.2    setosa
80         5.70        NA        3.5       1.0 versicolor
97         5.95        NA        4.2       1.3 versicolor
106        7.60        NA        6.6       2.1 virginica
```

# Now we have removed 1 NA's. Only 8 left

```
> iris_copy[is.na(iris_copy$Sepal.Width) & (iris_copy$Species == "setosa"), "Sepal.Width"] <- median(iris_copy$Sepal.Width[which(iris_copy$Species == "setosa")], na.rm = TRUE)
> iris_NA <- iris_copy[!complete.cases(iris_copy), ]
> iris_NA
  Sepal.Length Sepal.Width Petal.Length Petal.Width Species
5          5.00     3.6         NA       0.2   setosa
17         5.40     3.9         NA       0.4   setosa
35         4.90     3.1         NA       0.2   setosa
49         5.30     3.7         NA       0.2   setosa
80         5.70     NA         3.5      1.0 versicolor
97         5.95     NA         4.2      1.3 versicolor
```

# Now we have removed 1 NA's. Only 7 left

```
> iris_copy[is.na(iris_copy$Petal.Length) & (iris_copy$Species == "setosa"), "Petal.Length"] <- median(iris_copy$Petal.Length[which(iris_copy$Species == "setosa")], na.rm = TRUE)
> iris_NA <- iris_copy[!complete.cases(iris_copy), ]
> iris_NA
  Sepal.Length Sepal.Width Petal.Length Petal.Width Species
80          5.70     NA         3.5      1.0 versicolor
97          5.95     NA         4.2      1.3 versicolor
106         7.60     NA         6.6      2.1 virginica
```

*# Now it's your time to finish the logic and remove all the remaining NA's of the table*

*# We have saved a lot of interesting data. However we must be careful because each case is different and we can affect the result if we introduce the wrong number in the NA position*

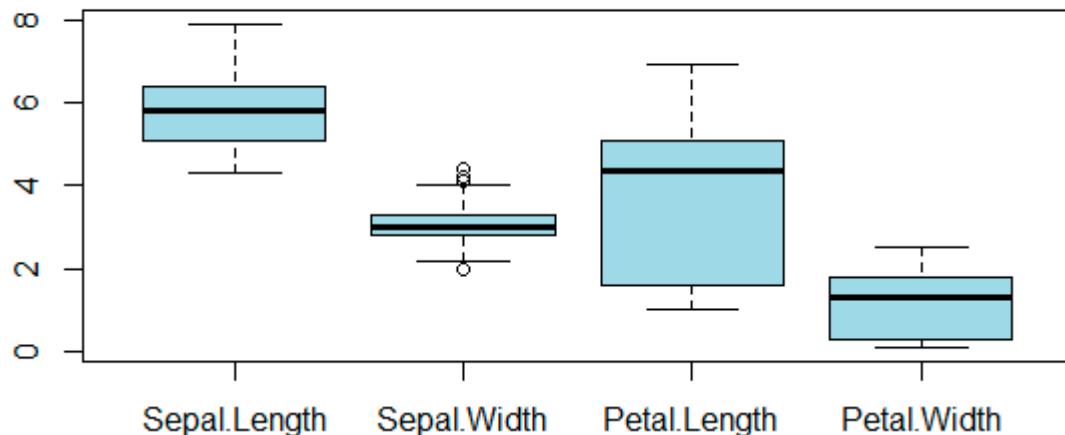
## B) Detecting Outliers

### Boxplot Method

Boxplots are useful for identifying outliers. In a boxplot, data points that lie beyond 1.5 times the interquartile range (IQR) are considered outliers.

```
> boxplot(iris[, 1:4], main = "Boxplot for Iris Features", col = "lightblue")
```

### Boxplot for Iris Features



- If any points are plotted outside the whiskers, they are considered **outliers**.

To find specific outliers:

```
> detect_outliers <- function(x) {
+   Q1 <- quantile(x, 0.25) # First quartile (25%)
+   Q3 <- quantile(x, 0.75) # Third quartile (75%)
+   IQR_value <- Q3 - Q1 # Interquartile range
+
+   lower_bound <- Q1 - 1.5 * IQR_value
+   upper_bound <- Q3 + 1.5 * IQR_value
+
+   return(which(x < lower_bound | x > upper_bound))
+ }
> outliers <- lapply(iris[, 1:4], detect_outliers)
> outliers
```

```
$Sepal.Length
integer(0)

$Sepal.Width
[1] 16 33 34 61

$Petal.Length
integer(0)

$Petal.Width
integer(0)

> |
```

#### 4. Visualization Using Scatterplots

*Scatterplots can help visualize potential outliers.*

```
r
CopyEdit
# Scatter plot matrix
pairs(iris[, 1:4], col = ifelse(rownames(iris) %in% unlist(outliers),
"red", "black"),
main = "Scatterplot Matrix with Outliers (Red)")
```

---

#### 5. Outlier Detection Using *outliers* Package

*You can also use the *outliers* package:*

```
r
CopyEdit
install.packages("outliers")
library(outliers)

# Identify the most extreme value in Sepal.Length
outlier_value <- scores(iris$Sepal.Length, type = "z")
outliers_sepals <- which(abs(outlier_value) > 3)
outliers_sepals
```

### C) Perform One-Hot Encoding

One-hot encoding is used to convert categorical variables into numerical format for machine learning models. In the **Iris dataset**, the `Species` column is categorical and can be converted into numerical dummy variables.

### Use `model.matrix()` for One-Hot Encoding

The `model.matrix()` function in R can be used to perform one-hot encoding:

```
> one_hot_iris <- model.matrix(~Species - 1, data = iris)
> head(one_hot_iris)
  Speciessetosa Speciesversicolor speciesvirginica
1          1                  0                 0
2          1                  0                 0
3          1                  0                 0
4          1                  0                 0
5          1                  0                 0
6          1                  0                 0
>
```

~Species - 1: The -1 removes the intercept term to ensure we get separate columns for each category.

This will create three new columns: `Speciessetosa`, `Speciesversicolor`, and `Speciesvirginica`, each containing 0 or 1.

### 3. Combine One-Hot Encoding with Original Data

To create a new dataframe with encoded values:

```
r
Copy
Edit
# Convert to dataframe
one_hot_iris <- as.data.frame(one_hot_iris)

# Combine with original numerical features
iris_encoded <- cbind(iris[, 1:4], one_hot_iris)

# View the updated dataset
head(iris_encoded)
```

### 4. Using caret Package for One-Hot Encoding

Alternatively, you can use the `caret` package:

```
r
Copy
Edit
install.packages("caret") # Install caret package if not installed
library(caret)
```

```
# Apply one-hot encoding
dummy_vars <- dummyVars(~ Species, data = iris)
iris_transformed <- predict(dummy_vars, newdata = iris)

# Convert to dataframe
iris_encoded <- cbind(iris[, 1:4], as.data.frame(iris_transformed))

# View result
head(iris_encoded)
```

#### D) Feature scaling (min-max and z-score normalization)

### Min-Max Scaling

Min-Max Scaling transforms the data to a fixed range, typically  $[0, 1]$ . The formula is:

$$X' = \frac{X - X_{\min}}{X_{\max} - X_{\min}}$$

```
> min_max_scaling <- function(x) {
+   return((x - min(x)) / (max(x) - min(x)))
+ }
>
> # Apply to numerical columns
> iris_minmax <- as.data.frame(lapply(iris[, 1:4], min_max_scaling))
>
> # Combine with Species column
> iris_minmax$Species <- iris$Species
>
> # View the transformed dataset
> head(iris_minmax)
```

```
Sepal.Length Sepal.Width Petal.Length Petal.Width Species
1  0.22222222  0.6250000  0.06779661  0.04166667 setosa
2  0.16666667  0.4166667  0.06779661  0.04166667 setosa
3  0.11111111  0.5000000  0.05084746  0.04166667 setosa
4  0.08333333  0.4583333  0.08474576  0.04166667 setosa
5  0.19444444  0.6666667  0.06779661  0.04166667 setosa
6  0.30555556  0.7916667  0.11864407  0.12500000 setosa
> |
```

## Z-Score Normalization (Standardization)

Z-Score Normalization transforms data to have **mean = 0** and **standard deviation = 1** using the formula:

$$X' = \frac{X - \mu}{\sigma}$$

where:

- $\mu$  = mean of the feature
- $\sigma$  = standard deviation of the feature

```
> # Function for z-score Normalization
> z_score_scaling <- function(x) {
+   return((x - mean(x)) / sd(x))
+ }
>
> # Apply to numerical columns
> iris_zscore <- as.data.frame(lapply(iris[, 1:4], z_score_scaling))
>
> # Combine with species column
> iris_zscore$Species <- iris$Species
>
> # View the transformed dataset
> head(iris_zscore)
  Sepal.Length Sepal.Width Petal.Length Petal.Width Species
1  -0.8976739  1.01560199   -1.335752   -1.311052 setosa
2  -1.1392005 -0.13153881   -1.335752   -1.311052 setosa
```

### 4. Using caret Package for Feature Scaling

You can also use the `caret` package:

```
r
CopyEdit
install.packages("caret") # Install caret package if not installed
library(caret)

# Min-Max Scaling
preprocess_minmax <- preProcess(iris[, 1:4], method = "range")
iris_minmax_caret <- predict(preprocess_minmax, iris[, 1:4])

# Z-Score Normalization
preprocess_zscore <- preProcess(iris[, 1:4], method = "center", "scale")
iris_zscore_caret <- predict(preprocess_zscore, iris[, 1:4])
```

```
# Add Species column back
iris_minmax_caret$Species <- iris$Species
iris_zscore_caret$Species <- iris$Species

# View results
head(iris_minmax_caret)
head(iris_zscore_caret)
```

### E) Generate polynomial features.

#### Using `poly()` Function

The `poly()` function generates orthogonal polynomials of a given degree.

```
> # Generate second-degree polynomial features for Sepal.Length
> iris$Sepal.Length_poly2 <- poly(iris$Sepal.Length, degree = 2, raw = TRUE)[, 2]
>
> # Generate third-degree polynomial features for all numerical columns
> iris_poly <- as.data.frame(lapply(iris[, 1:4], function(x) poly(x, degree = 3, raw = TRUE)))
>
> # Rename columns for clarity
> colnames(iris_poly) <- c("Sepal.Length_1", "Sepal.Length_2", "Sepal.Length_3",
+                           "Sepal.Width_1", "Sepal.Width_2", "Sepal.Width_3",
+                           "Petal.Length_1", "Petal.Length_2", "Petal.Length_3",
+                           "Petal.Width_1", "Petal.Width_2", "Petal.Width_3")
>
> # Add the Species column back
> iris_poly$Species <- iris$Species
>
> # View transformed dataset
> head(iris_poly)
```

- `poly(x, degree = 3, raw = TRUE)`: Creates polynomial features up to **degree 3**.
- `raw = TRUE`: Ensures regular polynomial terms (without orthogonal transformation).
  - `lapply()`: Applies the function to all numeric columns.

```
Using caret Package
The caret package allows automatic polynomial feature generation.
r
CopyEdit
install.packages("caret") # Install if not installed
library(caret)

# Generate polynomial features up to degree 3
poly_features <- preProcess(iris[, 1:4], method = "poly", degree = 3, raw = TRUE)

# Apply transformation
iris_poly_caret <- predict(poly_features, iris[, 1:4])
```

```
# Add Species column back
iris_poly_caret$Species <- iris$Species

# View results
head(iris_poly_caret)



---


4. Using polyreg Package
The polyreg package provides another way to generate polynomial features.
r
CopyEdit
install.packages("polyreg") # Install package
library(polyreg)

# Generate polynomial features up to degree 3
iris_polyreg <- poly_fit(iris[, 1:4], degree = 3)

# View transformed dataset
head(iris_polyreg)

Students have to explore the datasets within R libraries, Kaggle, UCI or any other freely available data repositories.
```

```
> # Load dataset
> data()
> str(airquality)
'data.frame': 153 obs. of 6 variables:
 $ Ozone : int 41 36 12 18 NA 28 23 19 8 NA ...
 $ Solar.R: int 190 118 149 313 NA NA 299 99 19 194 ...
 $ Wind : num 7.4 8 12.6 11.5 14.3 14.9 8.6 13.8 20.1 8.6 ...
 $ Temp : int 67 72 74 62 56 66 65 59 61 69 ...
 $ Month : int 5 5 5 5 5 5 5 5 5 ...
 $ Day : int 1 2 3 4 5 6 7 8 9 10 ...

>
> # Create a copy
> airquality_copy <- airquality
>
> # Introduce NA values manually
> airquality_copy$Ozone[c(5, 15, 30, 50, 75)] <- NA
> airquality_copy$Solar.R[c(10, 20, 40, 60, 90)] <- NA
> airquality_copy$Wind[c(7, 25, 55, 85)] <- NA
> airquality_copy$Temp[c(12, 35, 70)] <- NA
>
> # Summary of the dataset
> summary(airquality_copy)
   Ozone          Solar.R          Wind          Temp      
Min.   : 1.00   Min.   : 7.0   Min.   : 1.700   Min.   :56.00  
1st Qu.:18.00   1st Qu.:118.0  1st Qu.: 7.400   1st Qu.:72.25  
Median :32.00   Median :207.0  Median : 9.700   Median :79.00  
Mean   :41.96   Mean   :186.6  Mean   : 9.956   Mean   :77.81  
2nd Qu.:63.00   2nd Qu.:125.0  2nd Qu.:11.500   2nd Qu.:84.75   
```

3rd Qu.: 63.00	3rd Qu.: 258.0	3rd Qu.: 11.500	3rd Qu.: 84.75
Max. : 168.00	Max. : 334.0	Max. : 20.700	Max. : 97.00
NA's : 40	NA's : 12	NA's : 4	NA's : 3
Month	Day		
Min. : 5.000	Min. : 1.0		
1st Qu.: 6.000	1st Qu.: 8.0		
Median : 7.000	Median : 16.0		
Mean : 6.993	Mean : 15.8		
3rd Qu.: 8.000	3rd Qu.: 23.0		
Max. : 9.000	Max. : 31.0		

```
>
> # Count NA values
> length(which(is.na(airquality_copy)))
[1] 59
>
> # Extract rows with NA values
> airquality_NA <- airquality_copy[!complete.cases(airquality_copy),]
> airquality_NA
   Ozone Solar.R Wind Temp Month Day
5     NA      NA 14.3   56     5    5
6     28      NA 14.9   66     5    6
7     23     299    NA  65     5    7
10    NA      NA  8.6   69     5   10
11     7      NA  6.9   74     5   11
12     16     256  9.7    NA     5   12
15    NA      65 13.2   58     5   15
34    11     242 18.1   67     6   23
35    NA     186  9.2    NA     6    4
36    NA     220  8.6   85     6    5
37    NA     264 14.3   79     6    6
39    NA     273  6.9   87     6    8
40    71      NA 13.8   90     6    9
42    NA     259 10.9   93     6   11
43    NA     250  9.2   92     6   12
45    NA     332 13.8   80     6   14
46    NA     322 11.5   79     6   15
50    NA     120 11.5   73     6   19
52    NA     150  6.3   77     6   21
53    NA      59  1.7   76     6   22
54    NA      91  4.6   76     6   23
55    NA     250    NA  76     6   24
56    NA     135  8.0   75     6   25
57    NA     127  8.0   78     6   26
58    NA      47 10.3   73     6   27
59    NA      98 11.5   80     6   28
60    NA      NA 14.9   77     6   29
61    NA     138  8.0   83     6   30
65    NA     101 10.9   84     7    4
70     97     272  5.7    NA     7    9
72    NA     139  8.6   82     7   11
75    NA     291 14.9   91     7   14
83    NA     258  9.7   81     7   22
84    NA     295 11.5   82     7   23
85     80     294    NA  86     7   24
```

90	50	NA	7.4	86	7	29
96	78	NA	6.9	86	8	4
97	35	NA	7.4	85	8	5
98	66	NA	4.6	87	8	6
102	NA	222	8.6	92	8	10
103	NA	137	11.5	86	8	11
107	NA	64	11.5	79	8	15
115	NA	255	12.6	75	8	23
119	NA	153	5.7	88	8	27
150	NA	145	13.2	77	9	27
>						
> airquality_NA <- airquality_copy[rowSums(is.na(airquality_copy)) > 0,]						
> airquality_NA						
Ozone	Solar.R	Wind	Temp	Month	Day	
5	NA	NA	14.3	56	5	5
6	28	NA	14.9	66	5	6
7	23	299	NA	65	5	7
10	NA	NA	8.6	69	5	10
11	7	NA	6.9	74	5	11
12	16	256	9.7	NA	5	12
15	NA	65	13.2	58	5	15
20	11	NA	9.7	62	5	20
25	NA	66	NA	57	5	25
26	NA	266	14.9	58	5	26
27	NA	NA	8.0	57	5	27
30	NA	223	5.7	79	5	30
32	NA	286	8.6	78	6	1
33	NA	207	8.7	74	6	2
36	NA	220	8.6	85	6	5
37	NA	264	14.3	79	6	6
39	NA	273	6.9	87	6	8
40	71	NA	13.8	90	6	9
42	NA	259	10.9	93	6	11
43	NA	250	9.2	92	6	12
45	NA	332	13.8	80	6	14
46	NA	322	11.5	79	6	15
50	NA	120	11.5	73	6	19
52	NA	150	6.3	77	6	21
53	NA	59	1.7	76	6	22
54	NA	91	4.6	76	6	23
55	NA	250	NA	76	6	24
56	NA	135	8.0	75	6	25
57	NA	127	8.0	78	6	26
58	NA	47	10.3	73	6	27
59	NA	98	11.5	80	6	28
60	NA	NA	14.9	77	6	29
61	NA	138	8.0	83	6	30
65	NA	101	10.9	84	7	4
70	97	272	5.7	NA	7	9
72	NA	139	8.6	82	7	11
75	NA	291	14.9	91	7	14
83	NA	258	9.7	81	7	22
84	NA	295	11.5	82	7	23
85	80	294	NA	86	7	24
90	50	NA	7.4	86	7	29
--	--	--	--	--	-	-

```

119    NA    153  5.7   88     8  27
150    NA    145 13.2   77     9  27
>
> # Identify complete cases
> complete.cases(airquality_copy)
[1] TRUE TRUE TRUE TRUE FALSE FALSE FALSE TRUE TRUE TRUE FALSE FALSE FALSE
[13] TRUE TRUE FALSE TRUE TRUE TRUE FALSE TRUE TRUE TRUE TRUE TRUE TRUE
[25] FALSE FALSE FALSE TRUE TRUE FALSE TRUE FALSE FALSE FALSE FALSE FALSE FALSE
[37] FALSE TRUE FALSE FALSE TRUE FALSE FALSE FALSE TRUE FALSE FALSE TRUE TRUE
[49] TRUE FALSE TRUE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
[61] FALSE TRUE TRUE TRUE FALSE TRUE TRUE TRUE TRUE TRUE FALSE TRUE FALSE
[73] TRUE TRUE FALSE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE FALSE FALSE
[85] FALSE TRUE TRUE TRUE TRUE FALSE TRUE TRUE TRUE TRUE TRUE TRUE FALSE
[97] FALSE FALSE TRUE TRUE TRUE TRUE FALSE FALSE TRUE TRUE TRUE FALSE TRUE
[109] TRUE TRUE TRUE TRUE TRUE TRUE FALSE TRUE TRUE TRUE TRUE FALSE TRUE
[121] TRUE TRUE
[133] TRUE TRUE
[145] TRUE TRUE TRUE TRUE TRUE FALSE TRUE TRUE TRUE TRUE TRUE TRUE TRUE
>
> # Create a clean dataset without NA values
> airquality_clean <- airquality_copy[complete.cases(airquality_copy),]
> length(which(is.na(airquality_clean)))
[1] 0
>
> # Impute missing values with median based on Month
> airquality_copy[is.na(airquality_copy$Ozone) & (airquality_copy$Month == 5), "Ozone"] <- median(airquality_copy$Ozone[which(airquality_copy$Month == 5)], na.rm = TRUE)

```

	Ozone	Solar.R	Wind	Temp	Month	Day
5	17	NA	14.3	56	5	5
6	28	NA	14.9	66	5	6
7	23	299	NA	65	5	7
10	17	NA	8.6	69	5	10
11	7	NA	6.9	74	5	11
12	16	256	9.7	NA	5	12
20	11	NA	9.7	62	5	20
25	17	66	NA	57	5	25
27	17	NA	8.0	57	5	27
32	NA	286	8.6	78	6	1
33	NA	287	9.7	74	6	2
34	NA	242	16.1	67	6	3
35	NA	186	9.2	NA	6	4
36	NA	220	8.6	85	6	5
37	NA	264	14.3	79	6	6
39	NA	273	6.9	87	6	8
40	71	NA	13.8	90	6	9
42	NA	259	10.9	93	6	11
43	NA	250	9.2	92	6	12
45	NA	332	13.8	80	6	14
46	NA	322	11.5	79	6	15
50	NA	120	11.5	73	6	19
52	NA	150	6.3	77	6	21
53	NA	59	1.7	76	6	22
54	NA	91	4.6	76	6	23
55	NA	250	NA	76	6	24
..	..	..	..	..	..	..

```

> airquality_copy[is.na(airquality_copy$Solar.R) & (airquality_copy$Month == 6), "Solar.R"] <- median(airquality_copy$Solar.R[which(airquality_copy$Month == 6)], na.rm = TRUE)
> airquality_NA <- airquality_copy[!complete.cases(airquality_copy),]
> airquality_NA
   Ozone Solar.R Wind Temp Month Day
5      17      NA 14.3  56     5    5
6      28      NA 14.9  66     5    6
7      23    299.0   NA  65     5    7
10     17      NA  8.6  69     5   10
11      7      NA  6.9  74     5   11
12     16   256.0  9.7  NA     5   12
20     11      NA  9.7  62     5   20
25     17    66.0   NA  57     5   25
27     17      NA  8.0  57     5   27
32     NA   286.0  8.6  78     6    1
33     NA   287.0  9.7  74     6    2
34     NA   242.0 16.1  67     6    3
35     NA   186.0  9.2  NA     6    4
36     NA   220.0  8.6  85     6    5
37     NA   264.0 14.3  79     6    6
39     NA   273.0  6.9  87     6    8
42     NA   259.0 10.9  93     6   11
43     NA   250.0  9.2  92     6   12
45     NA   332.0 13.8  80     6   14
46     NA   322.0 11.5  79     6   15
50     NA   120.0 11.5  73     6   19
52     NA   150.0  6.3  77     6   21

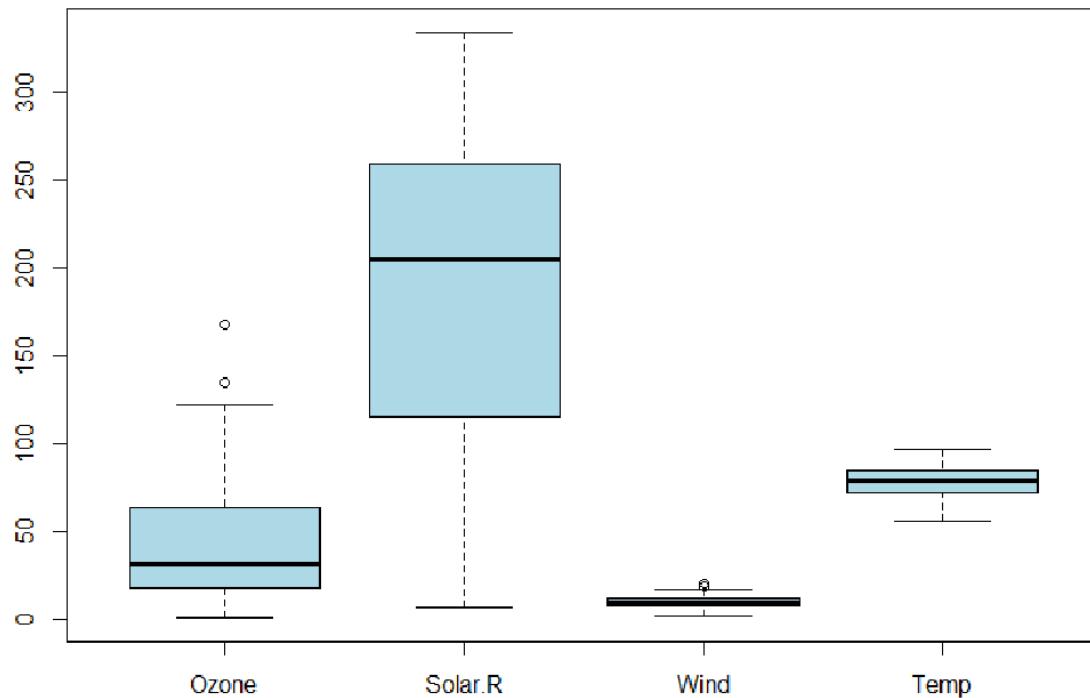
> airquality_copy[is.na(airquality_copy$Wind) & (airquality_copy$Month == 7), "Wind"]
<- median(airquality_copy$Wind[which(airquality_copy$Month == 7)], na.rm = TRUE)
> airquality_NA <- airquality_copy[!complete.cases(airquality_copy),]
> airquality_NA
   Ozone Solar.R Wind Temp Month Day
5      17      NA 14.3  56     5    5
6      28      NA 14.9  66     5    6
7      23    299.0   NA  65     5    7
10     17      NA  8.6  69     5   10
11      7      NA  6.9  74     5   11
12     16   256.0  9.7  NA     5   12
20     11      NA  9.7  62     5   20
25     17    66.0   NA  57     5   25
27     17      NA  8.0  57     5   27
32     NA   286.0  8.6  78     6    1
33     NA   287.0  9.7  74     6    2
34     NA   242.0 16.1  67     6    3
35     NA   186.0  9.2  NA     6    4
36     NA   220.0  8.6  85     6    5
37     NA   264.0 14.3  79     6    6
39     NA   273.0  6.9  87     6    8
42     NA   259.0 10.9  93     6   11
43     NA   250.0  9.2  92     6   12
45     NA   332.0 13.8  80     6   14
46     NA   322.0 11.5  79     6   15
50     NA   120.0 11.5  73     6   19
52     NA   150.0  6.3  77     6   21

```

```

> airquality_copy[is.na(airquality_copy$Temp) & (airquality_copy$Month == 8), "Temp"]
<- median(airquality_copy$Temp[which(airquality_copy$Month == 8)], na.rm = TRUE)
> airquality_NA <- airquality_copy[!complete.cases(airquality_copy),]
> airquality_NA
   Ozone Solar.R Wind Temp Month Day
 5     17      NA 14.3  56     5    5
 6     28      NA 14.9  66     5    6
 7     23  299.0    NA  65     5    7
 10    17      NA  8.6  69     5   10
 11     7      NA  6.9  74     5   11
 12    16  256.0   9.7  NA     5   12
 20    11      NA  9.7  62     5   20
 25    17   66.0    NA  57     5   25
 27    17      NA  8.0  57     5   27
 32    NA  286.0   8.6  78     6    1
 33    NA  287.0   9.7  74     6    2
 34    NA  242.0  16.1  67     6    3
 35    NA  186.0   9.2  NA     6    4
 36    NA  220.0   8.6  85     6    5
 37    NA  264.0  14.3  79     6    6
 39    NA  273.0   6.9  87     6    8
 42    NA  259.0  10.9  93     6   11
 43    NA  250.0   9.2  92     6   12
 45    NA  332.0  13.8  80     6   14
 46    NA  322.0  11.5  79     6   15
 50    NA  120.0  11.5  73     6   19
 52    NA  150.0  10.2  77     6   21
  
```

**Boxplot for AirQuality Features**



```

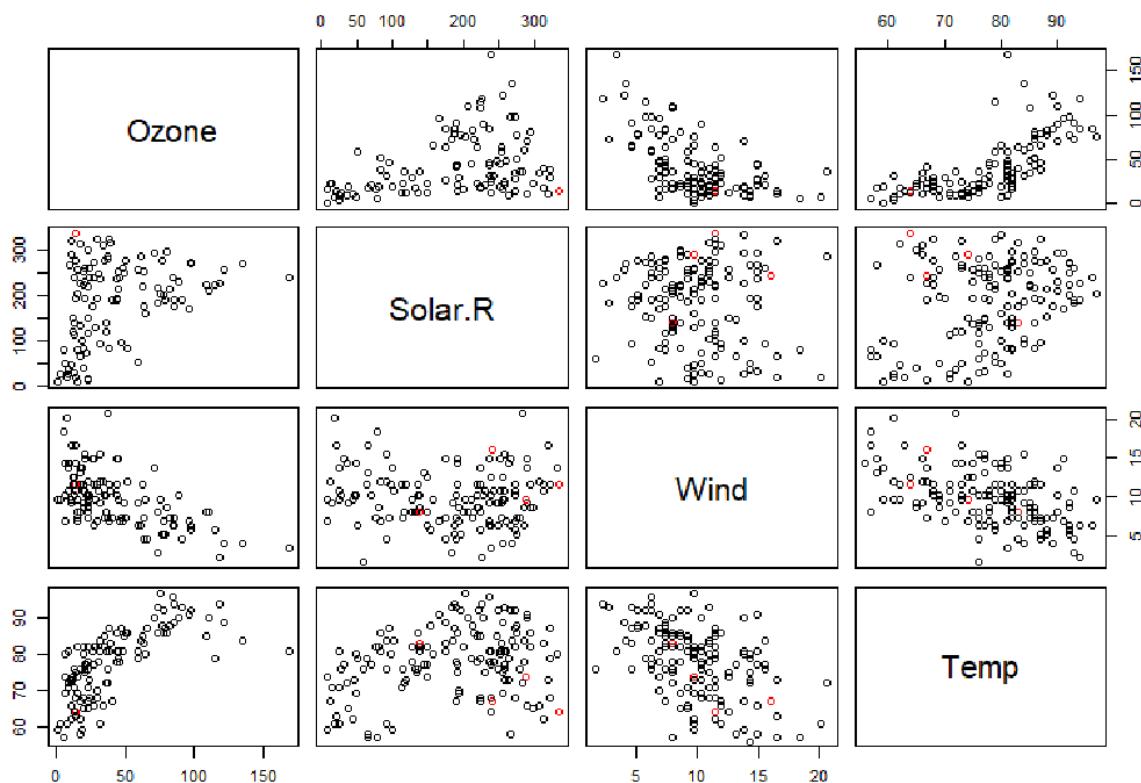
> # Detecting Outliers using Boxplot Method
> boxplot(airquality[, 1:4], main = "Boxplot for AirQuality Features", col = "lightblue")
>
> detect_outliers <- function(x) {
+   Q1 <- quantile(x, 0.25, na.rm = TRUE) # First quartile (25%)
+   Q3 <- quantile(x, 0.75, na.rm = TRUE) # Third quartile (75%)
+   IQR_value <- Q3 - Q1 # Interquartile range
+
+   lower_bound <- Q1 - 1.5 * IQR_value
+   upper_bound <- Q3 + 1.5 * IQR_value
+
+   return(which(x < lower_bound | x > upper_bound))
+ }
> outliers <- lapply(airquality[, 1:4], detect_outliers)
> outliers
$Ozone
[1] 62 117

$Solar.R
integer(0)

$Wind
[1] 9 18 48

$Temp
integer(0)
  
```

**Scatterplot Matrix with Outliers (Red)**



```

> # Scatterplot Matrix with Outliers Highlighted
> pairs(airquality[, 1:4], col = ifelse(rownames(airquality) %in% unlist(outliers),
  "red", "black"), main = "Scatterplot Matrix with Outliers (Red)")
>
> # Install and Load Outliers Package
> library(outliers)
>
> outlier_value <- scores(airquality$ozone, type = "z")
> outliers_ozone <- which(abs(outlier_value) > 3)
> outliers_ozone
integer(0)
>
> # Perform One-Hot Encoding
> one_hot_airquality <- model.matrix(~Month - 1, data = airquality)
> head(one_hot_airquality)
  Month
1      5
2      5
3      5
4      5
5      5
6      5
>
> # Convert to dataframe
> one_hot_airquality <- as.data.frame(one_hot_airquality)
>
> # Combine with original numerical features
> airquality_encoded <- cbind(airquality[, 1:4], one_hot_airquality)
>
> # View the updated dataset
> head(airquality_encoded)
  Ozone Solar.R Wind Temp Month
1    41     190   7.4   67     5
2    36     118   8.0   72     5
3    12     149  12.6   74     5
4    18     313  11.5   62     5
5    NA      NA  14.3   56     5
6    28     NA  14.9   66     5
>
> # Install and Load caret Package
> library(caret)
> |

```

## Red Colour part (commands) student should try

### Conclusion:

Learnt various techniques for data cleaning, transformation, and feature engineering. These included handling missing values, detecting outliers, performing one-hot encoding, label encoding, feature scaling (min-max and z-score normalization), and generating polynomial features. By applying these techniques to the Iris dataset, students learned how to prepare data for machine learning models, ensuring better model performance and insights.

### **Post Lab Question**

1) Discuss an alternative approach to capturing non-linear relationships in a dataset instead of polynomial feature generation.

- Kernel Methods: Use SVMs with non-linear kernels (e.g., RBF) to implicitly map data to higher dimensions.
- Neural Networks: Use layers with non-linear activation functions (e.g., ReLU) for complex patterns.
- Local Regression (LOESS): Fit simple models to localized data subsets.
- Bayesian Models: Use Gaussian Processes for non-linear relationships with uncertainty estimates.
- Rule-Based Systems: Build interpretable non-linear models with decision rules.

2) What is the difference between univariate and multivariate outlier detection?

Univariate Outlier Detection:

- Focuses on a single feature at a time.
- Uses statistical methods like Z-scores, IQR (Interquartile Range), or boxplots to identify extreme values.
- Ignores relationships between features, so it may miss outliers that are only apparent when considering multiple features together.

Multivariate Outlier Detection:

- Considers multiple features simultaneously.
- Captures relationships and interactions between features.
- Better suited for detecting complex outliers that depend on the interaction of multiple variables.

Key Difference: Univariate is simpler and feature-specific, while multivariate is more comprehensive and accounts for feature relationships.