



Chapter 4: Advanced SQL

Database System Concepts, 5th Ed.

©Silberschatz, Korth and Sudarshan
See www.db-book.com for conditions on re-use





Chapter 4: Advanced SQL

- SQL Data Types and Schemas
- Integrity Constraints
- Authorization
- Embedded SQL
- Dynamic SQL
- Functions and Procedural Constructs**
- Recursive Queries**
- Advanced SQL Features**





Built-in Data Types in SQL

- **date:** Dates, containing a (4 digit) year, month and date
 - Example: **date** '2005-7-27'
- **time:** Time of day, in hours, minutes and seconds.
 - Example: **time** '09:00:30' **time** '09:00:30.75'
- **timestamp:** date plus time of day
 - Example: **timestamp** '2005-7-27 09:00:30.75'
- **interval:** period of time
 - Example: **interval** '1' day
 - Subtracting a date/time/timestamp value from another gives an interval value
 - Interval values can be added to date/time/timestamp values





Build-in Data Types in SQL (Cont.)

- Can extract values of individual fields from date/time/timestamp
 - Example: **extract (year from rstarttime)**
- Can cast string types to date/time/timestamp
 - Example: **cast <string-valued-expression> as date**
 - Example: **cast <string-valued-expression> as time**





User-Defined Types

- **create type** construct in SQL creates user-defined type

```
create type Dollars as numeric (12,2) final
```

- **create domain** construct in SQL-92 creates user-defined domain types

```
create domain person_name char(20) not null
```

- Types and domains are similar. Domains can have constraints, such as **not null**, specified on them.





Domain Constraints

- **Domain constraints** are the most elementary form of integrity constraint. They test values inserted in the database, and test queries to ensure that the comparisons make sense.
- New domains can be created from existing data types
 - Example: **create domain Dollars numeric(12, 2)**
create domain Pounds numeric(12,2)
- We cannot assign or compare a value of type Dollars to a value of type Pounds.
 - However, we can convert type as below
(cast r.A as Pounds)
(Should also multiply by the dollar-to-pound conversion-rate)





Large-Object Types

- Large objects (photos, videos, CAD files, etc.) are stored as a *large object*:
 - **blob**: binary large object -- object is a large collection of uninterpreted binary data (whose interpretation is left to an application outside of the database system)
 - **clob**: character large object -- object is a large collection of character data
 - When a query returns a large object, a pointer is returned rather than the large object itself.





Integrity Constraints

- Integrity constraints guard against accidental damage to the database, by ensuring that authorized changes to the database do not result in a loss of data consistency.
 - A checking account must have a balance greater than \$10,000.00
 - A salary of a bank employee must be at least \$4.00 an hour
 - A customer must have a (non-null) phone number





Constraints on a Single Relation

- **not null**
- **primary key**
- **unique**
- **check** (P), where P is a predicate





Not Null Constraint

- Declare *branch_name* for *branch* is **not null**
branch_name **char(15) not null**
- Declare the domain *Dollars* to be **not null**

create domain Dollars numeric(12,2) not null





The Unique Constraint

- **unique** (A_1, A_2, \dots, A_m)
- The unique specification states that the attributes A_1, A_2, \dots, A_m form a candidate key.
- Candidate keys are permitted to be null (in contrast to primary keys).





The check clause

- **check (P)**, where P is a predicate

Example: Declare *branch_name* as the primary key for *branch* and ensure that the values of *assets* are non-negative.

```
create table branch
  (branch_name    char(15),
   branch_city     char(30),
   assets          integer,
   primary key (branch_name),
   check (assets >= 0))
```





The check clause (Cont.)

- The **check** clause in SQL-92 permits domains to be restricted:
 - Use **check** clause to ensure that an hourly_wage domain allows only values greater than a specified value.

```
create domain hourly_wage numeric(5,2)
constraint value_test check(hourly_wage > = 4.00)
```
 - The domain has a constraint that ensures that the hourly_wage is greater than 4.00
 - The clause **constraint** *value_test* is optional; useful to indicate which constraint an update violated.





Referential Integrity

- Ensures that a value that appears in one relation for a given set of attributes also appears for a certain set of attributes in another relation.
 - Example: If “Perryridge” is a branch name appearing in one of the tuples in the *account* relation, then there exists a tuple in the *branch* relation for branch “Perryridge”.
- Primary and candidate keys and foreign keys can be specified as part of the SQL **create table** statement:
 - The **primary key** clause lists attributes that comprise the primary key.
 - The **unique key** clause lists attributes that comprise a candidate key.
 - The **foreign key** clause lists the attributes that comprise the foreign key and the name of the relation referenced by the foreign key. By default, a foreign key references the primary key attributes of the referenced table.





Referential Integrity in SQL – Example

```
create table customer
  (customer_name      char(20),
   customer_street    char(30),
   customer_city      char(30),
   primary key (customer_name))

create table branch
  (branch_name char(15),
   branch_city  char(30),
   assets       numeric(12,2),
   primary key (branch_name))
```





Referential Integrity in SQL – Example (Cont.)

```
create table account
```

```
(account_number      char(10),  
branch_name   char(15),  
balance        integer,  
primary key (account_number),  
foreign key (branch_name) references branch )
```

```
create table depositor
```

```
(customer_name      char(20),  
account_number      char(10),  
primary key (customer_name, account_number),  
foreign key (account_number ) references account,  
foreign key (customer_name ) references customer )
```





Assertions

- An **assertion** is a predicate expressing a condition that we wish the database always to satisfy.
- An assertion in SQL takes the form
create assertion <assertion-name> check <predicate>
- When an assertion is made, the system tests it for validity, and tests it again on every update that may violate the assertion
 - This testing may introduce a significant amount of overhead; hence assertions should be used with great care.
- Asserting
for all X , $P(X)$
is achieved in a round-about fashion using
not exists X such that not $P(X)$





- **create assertion** <assertion-name> **check** <predicate>
- STUDENT(Roll no, name, Batch)

Create assertion Batch-size CHECK

```
(  
 ( SELECT BATCH , COUNT()  
   FROM STUDENT  
     GROUP BY BATCH HAVING COUNT <=20 )  
)
```





Assertion Example

- Every loan has at least one borrower who maintains an account with a minimum balance of \$1000.00

create assertion *balance_constraint* **check**

(not exists (

select *

from *loan*

where not exists (

select *

from *borrower, depositor, account*

where *loan.loan_number = borrower.loan_number*

and *borrower.customer_name =*

depositor.customer_name

and *depositor.account_number =*

account.account_number

and *account.balance >= 1000)))*





Assertion Example

- The sum of all loan amounts for each branch must be less than the sum of all account balances at the branch.

create assertion sum_constraint check

```
(not exists (select *
             from branch
             where (select sum(amount )
                    from loan
                    where loan.branch_name =
                          branch.branch_name )
                   >= (select sum (amount )
                        from account
                        where loan.branch_name =
                              branch.branch_name )))
```





Authorization

Forms of authorization on parts of the database:

- **Read** - allows reading, but not modification of data.
- **Insert** - allows insertion of new data, but not modification of existing data.
- **Update** - allows modification, but not deletion of data.
- **Delete** - allows deletion of data.

Forms of authorization to modify the database schema (covered in Chapter 8):

- **Index** - allows creation and deletion of indices.
- **Resources** - allows creation of new relations.
- **Alteration** - allows addition or deletion of attributes in a relation.
- **Drop** - allows deletion of relations.





Authorization Specification in SQL

- The **grant** statement is used to confer authorization
 - grant** <privilege list>
on <relation name or view name> **to** <user list>
- <user list> is:
 - a user-id
 - **public**, which allows all valid users the privilege granted
 - A role (more on this in Chapter 8)
- Granting a privilege on a view does not imply granting any privileges on the underlying relations.
- The grantor of the privilege must already hold the privilege on the specified item (or be the database administrator).





Privileges in SQL

- **select**: allows read access to relation, or the ability to query using the view
 - Example: grant users U_1 , U_2 , and U_3 **select** authorization on the *branch* relation:
grant select on branch to U_1 , U_2 , U_3
- **insert**: the ability to insert tuples
- **update**: the ability to update using the SQL update statement
- **delete**: the ability to delete tuples.
- **all privileges**: used as a short form for all the allowable privileges
- more in Chapter 8





Revoking Authorization in SQL

- The **revoke** statement is used to revoke authorization.
revoke <privilege list>
 on <relation name or view name> **from** <user list>
- Example:
 revoke select on branch from U₁, U₂, U₃
- All privileges that depend on the privilege being revoked are also revoked.
- <privilege-list> may be **all** to revoke all privileges the revoker may hold.
- If the same privilege was granted twice to the same user by different grantees, the user may retain the privilege after the revocation.





Embedded SQL

- The SQL standard defines embeddings of SQL in a variety of programming languages such as C, Java, and Cobol.
 - A language to which SQL queries are embedded is referred to as a **host language**, and the SQL structures permitted in the host language comprise *embedded SQL*.
 - The basic form of these languages follows that of the System R embedding of SQL into PL/I.
 - **EXEC SQL** statement is used to identify embedded SQL request to the preprocessor

EXEC SQL <embedded SQL statement> END EXEC

Note: this varies by language (for example, the Java embedding uses
SQL { };)





- Specify the query in SQL and declare a *cursor* for it
- Eg: To find the customer names and their city

EXEC SQL

```
declare c cursor for  
select name, city  
from customer
```

END_EXEC





Example Query

- From within a host language, find the names and cities of customers with more than the variable **amount** dollars in some account.

- Specify the query in SQL and declare a *cursor* for it

EXEC SQL

```
declare c cursor for
select depositor.customer_name, customer_city
from depositor, customer, account
where depositor.customer_name = customer.customer_name
      and depositor.account_number = account.account_number
      and account.balance > :amount
END_EXEC
```





Embedded SQL (Cont.)

- The **open** statement causes the query to be evaluated

```
EXEC SQL open c END_EXEC
```
- The **fetch** statement causes the values of one tuple in the query result to be placed on host language variables.

```
EXEC SQL fetch c into :cn, :cc END_EXEC
```

Repeated calls to **fetch** get successive tuples in the query result
- A variable called SQLSTATE in the SQL communication area (SQLCA) gets set to '02000' to indicate no more data is available
- The **close** statement causes the database system to delete the temporary relation that holds the result of the query.

```
EXEC SQL close c END_EXEC
```

Note: above details vary with language. For example, the Java embedding defines Java iterators to step through result tuples.





Updates Through Cursors

- Can update tuples fetched by cursor by declaring that the cursor is for update

```
declare c cursor for
select *
from account
where branch_name = 'Perryridge'
for update
```

- To update tuple at the current location of cursor c

```
update account
set balance = balance + 100
where current of c
```





Dynamic SQL

- Allows programs to construct and submit SQL queries at run time.
- Example of the use of dynamic SQL from within a C program.

```
char * sqlprog = "update account
                  set balance = balance * 1.05
                  where account_number = ?"
EXEC SQL prepare dynprog from :sqlprog;
char account [10] = "A-101";
EXEC SQL execute dynprog using :account;
```

- The dynamic SQL program contains a ?, which is a place holder for a value that is provided when the SQL program is executed.





ODBC and JDBC

- API (application-program interface) for a program to interact with a database server
- Application makes calls to
 - Connect with the database server
 - Send SQL commands to the database server
 - Fetch tuples of result one-by-one into program variables
- ODBC (Open Database Connectivity) works with C, C++, C#, and Visual Basic
- JDBC (Java Database Connectivity) works with Java





ODBC

- Open DataBase Connectivity(ODBC) standard
 - standard for application program to communicate with a database server.
 - application program interface (API) to
 - 4 open a connection with a database,
 - 4 send queries and updates,
 - 4 get back results.
- Applications such as GUI, spreadsheets, etc. can use ODBC





ODBC (Cont.)

- Each database system supporting ODBC provides a "driver" library that must be linked with the client program.
- When client program makes an ODBC API call, the code in the library communicates with the server to carry out the requested action, and fetch results.
- ODBC program first allocates an SQL environment, then a database connection handle.
- Opens database connection using SQLConnect(). Parameters for SQLConnect:
 - connection handle,
 - the server to which to connect
 - the user identifier,
 - password
- Must also specify types of arguments:
 - SQL_NTS denotes previous argument is a null-terminated string.





ODBC Code

- ```
int ODBCexample()
{
 RETCODE error;
 HENV env; /* environment */
 HDBC conn; /* database connection */
 SQLAllocEnv(&env);
 SQLAllocConnect(env, &conn);
 SQLConnect(conn, "aura.bell-labs.com", SQL_NTS, "avi", SQL_NTS,
 "avipasswd", SQL_NTS);
 { Do actual work ... }

 SQLDisconnect(conn);
 SQLFreeConnect(conn);
 SQLFreeEnv(env);
}
```





# ODBC Code (Cont.)

- Program sends SQL commands to the database by using SQLExecDirect
- Result tuples are fetched using SQLFetch()
- SQLBindCol() binds C language variables to attributes of the query result
  - When a tuple is fetched, its attribute values are automatically stored in corresponding C variables.
  - Arguments to SQLBindCol()
    - 4 ODBC stmt variable, attribute position in query result
    - 4 The type conversion from SQL to C.
    - 4 The address of the variable.
    - 4 For variable-length types like character arrays,
      - The maximum length of the variable
      - Location to store actual length when a tuple is fetched.
      - Note: A negative value returned for the length field indicates null value
- Good programming requires checking results of every function call for errors; we have omitted most checks for brevity.





# ODBC Code (Cont.)

- Main body of program

```
char branchname[80];
float balance;
int lenOut1, lenOut2;
HSTMT stmt;
SQLAllocStmt(conn, &stmt);
char * sqlquery = "select branch_name, sum (balance)
 from account
 group by branch_name";
error = SQLExecDirect(stmt, sqlquery, SQL_NTS);
if (error == SQL_SUCCESS) {
 SQLBindCol(stmt, 1, SQL_C_CHAR, branchname , 80,
&lenOut1);
 SQLBindCol(stmt, 2, SQL_C_FLOAT, &balance, 0 ,
&lenOut2);
 while (SQLFetch(stmt) >= SQL_SUCCESS) {
 printf ("%s %g\n", branchname, balance);
 }
}
SQLFreeStmt(stmt, SQL_DROP);
```





# More ODBC Features

- **Prepared Statement**
  - SQL statement prepared: compiled at the database
  - Can have placeholders: E.g. insert into account values(?, ?, ?)
  - Repeatedly executed with actual values for the placeholders
- **Metadata features**
  - finding all the relations in the database and
  - finding the names and types of columns of a query result or a relation in the database.
- By default, each SQL statement is treated as a separate transaction that is committed automatically.
  - Can turn off automatic commit on a connection
    - 4 SQLSetConnectOption(conn, SQL\_AUTOCOMMIT, 0)}
  - transactions must then be committed or rolled back explicitly by
    - 4 SQLTransact(conn, SQL\_COMMIT) or
    - 4 SQLTransact(conn, SQL\_ROLLBACK)





# ODBC Conformance Levels

- Conformance levels specify subsets of the functionality defined by the standard.
  - Core
  - Level 1 requires support for metadata querying
  - Level 2 requires ability to send and retrieve arrays of parameter values and more detailed catalog information.
- SQL Call Level Interface (CLI) standard similar to ODBC interface, but with some minor differences.





# JDBC

- JDBC is a Java API for communicating with database systems supporting SQL
- JDBC supports a variety of features for querying and updating data, and for retrieving query results
- JDBC also supports metadata retrieval, such as querying about relations present in the database and the names and types of relation attributes
- Model for communicating with the database:
  - Open a connection
  - Create a “statement” object
  - Execute queries using the Statement object to send queries and fetch results
  - Exception mechanism to handle errors





# JDBC Code

```
public static void JDBCexample(String dbid, String userid, String passwd)
{
 try {
 Class.forName ("oracle.jdbc.driver.OracleDriver");
 Connection conn = DriverManager.getConnection(
 "jdbc:oracle:thin:@aura.bell-labs.com:2000:bankdb", userid, passwd);
 Statement stmt = conn.createStatement();
 ... Do Actual Work
 stmt.close();
 conn.close();
 }
 catch (SQLException sqle) {
 System.out.println("SQLException : " + sqle);
 }
}
```





# JDBC Code (Cont.)

- Update to database

```
try {
 stmt.executeUpdate("insert into account values
 ('A-9732', 'Perryridge', 1200)");
} catch (SQLException sqle) {
 System.out.println("Could not insert tuple. " + sqle);
}
```

- Execute query and fetch and print results

```
ResultSet rset = stmt.executeQuery("select branch_name,
 avg(balance)
 from account
 group by branch_name");

while (rset.next()) {
 System.out.println(
 rset.getString("branch_name") + " " + rset.getFloat(2));

}
```





# JDBC Code Details

- Getting result fields:
  - **rs.getString("branchname") and rs.getString(1) equivalent if branchname is the first argument of select result.**
- Dealing with Null values

```
int a = rs.getInt("a");
if (rs.wasNull()) System.out.println("Got null value");
```





# Procedural Extensions and Stored Procedures

- SQL provides a **module** language
  - Permits definition of procedures in SQL, with if-then-else statements, for and while loops, etc.
  - more in Chapter 9
- Stored Procedures
  - Can store procedures in the database
  - then execute them using the **call** statement
  - permit external applications to operate on the database without knowing about internal details
- These features are covered in Chapter 9 (Object Relational Databases)





# Functions and Procedures

- SQL:1999 supports functions and procedures
  - Functions/procedures can be written in SQL itself, or in an external programming language
  - Functions are particularly useful with specialized data types such as images and geometric objects
    - 4 Example: functions to check if polygons overlap, or to compare images for similarity
  - Some database systems support **table-valued functions**, which can return a relation as a result
- SQL:1999 also supports a rich set of imperative constructs, including
  - Loops, if-then-else, assignment
- Many databases have proprietary procedural extensions to SQL that differ from SQL:1999





# SQL Functions

- Define a function that, given the name of a customer, returns the count of the number of accounts owned by the customer.

```
create function account_count (customer_name varchar(20))
returns integer
begin
 declare a_count integer;
 select count (*) into a_count
 from depositor
 where depositor.customer_name = customer_name
 return a_count;
end
```

- Find the name and address of each customer that has more than one account.

```
select customer_name, customer_street, customer_city
from customer
where account_count (customer_name) > 1
```





# Table Functions

- SQL:2003 added functions that return a relation as a result
- Example: Return all accounts owned by a given customer

```
create function accounts_of (customer_name char(20)
```

```
 returns table (account_number char(10),
 branch_name char(15)
 balance numeric(12,2))
```

```
return table
```

```
(select account_number, branch_name, balance
 from account A
 where exists (
 select *
 from depositor D
 where D.customer_name = accounts_of.customer_name
 and D.account_number = A.account_number))
```





# Table Functions (cont'd)

- Usage

```
select *
from table (accounts_of ('Smith'))
```





# SQL Procedures

- The *author\_count* function could instead be written as procedure:

```
create procedure account_count_proc (in title varchar(20),
 out a_count integer)
begin
 select count(author) into a_count
 from depositor
 where depositor.customer_name =
account_count_proc.customer_name
end
```

- Procedures can be invoked either from an SQL procedure or from embedded SQL, using the **call** statement.

```
declare a_count integer;
call account_count_proc('Smith', a_count);
```

Procedures and functions can be invoked also from dynamic SQL

- SQL:1999 allows more than one function/procedure of the same name (called name **overloading**), as long as the number of arguments differ, or at least the types of the arguments differ





# Procedural Constructs

- Compound statement: **begin ... end**,
  - May contain multiple SQL statements between **begin** and **end**.
  - Local variables can be declared within a compound statements
- **While** and **repeat** statements:

```
declare n integer default 0;
```

```
while n < 10 do
```

```
 set n = n + 1
```

```
end while
```

```
repeat
```

```
 set n = n - 1
```

```
until n = 0
```

```
end repeat
```





# Procedural Constructs (Cont.)

- **For** loop
  - Permits iteration over all results of a query
  - Example: find total of all balances at the Perryridge branch

```
declare n integer default 0;
for r as
 select balance from account
 where branch_name = 'Perryridge'
do
 set n = n + r.balance
end for
```





# Procedural Constructs (cont.)

- Conditional statements (**if-then-else**)

E.g. To find sum of balances for each of three categories of accounts (with balance <1000, >=1000 and <5000, >= 5000)

```
if r.balance < 1000
 then set l = l + r.balance
elseif r.balance < 5000
 then set m = m + r.balance
else set h = h + r.balance
end if
```

- SQL:1999 also supports a **case** statement similar to C case statement
- Signaling of exception conditions, and declaring handlers for exceptions

```
declare out_of_stock condition
declare exit handler for out_of_stock
begin
 ...
 .. signal out-of-stock
end
```

- The handler here is **exit** -- causes enclosing **begin..end** to be exited
- Other actions possible on exception





# External Language Functions/Procedures

- SQL:1999 permits the use of functions and procedures written in other languages such as C or C++
- Declaring external language procedures and functions

```
create procedure account_count_proc(in customer_name varchar(20),
 out count integer)
```

**language** C

**external name** '/usr/avi/bin/account\_count\_proc'

```
create function account_count(customer_name varchar(20))
```

**returns** integer

**language** C

**external name** '/usr/avi/bin/author\_count'





# External Language Routines (Cont.)

- Benefits of external language functions/procedures:
  - more efficient for many operations, and more expressive power
- Drawbacks
  - Code to implement function may need to be loaded into database system and executed in the database system's address space
    - 4 risk of accidental corruption of database structures
    - 4 security risk, allowing users access to unauthorized data
  - There are alternatives, which give good security at the cost of potentially worse performance
  - Direct execution in the database system's space is used when efficiency is more important than security





# Security with External Language Routines

- To deal with security problems
  - Use **sandbox** techniques
    - 4 that is use a safe language like Java, which cannot be used to access/damage other parts of the database code
  - Or, run external language functions/procedures in a separate process, with no access to the database process' memory
    - 4 Parameters and results communicated via inter-process communication
- Both have performance overheads
- Many database systems support both above approaches as well as direct executing in database system address space





# Recursion in SQL

- SQL:1999 permits recursive view definition
- Example: find all employee-manager pairs, where the employee reports to the manager directly or indirectly (that is manager's manager, manager's manager's manager, etc.)

```
with recursive empl (employee_name, manager_name) as (
 select employee_name, manager_name
 from manager
 union
 select manager.employee_name, empl.manager_name
 from manager, empl
 where manager.manager_name = empl.employee_name)
select *
 from empl
```

This example view, *empl*, is called the *transitive closure* of the *manager* relation





# The Power of Recursion

- Recursive views make it possible to write queries, such as transitive closure queries, that cannot be written without recursion or iteration.
  - Intuition: Without recursion, a non-recursive non-iterative program can perform only a fixed number of joins of *manager* with itself
    - 4 This can give only a fixed number of levels of managers
    - 4 Given a program we can construct a database with a greater number of levels of managers on which the program will not work
- Computing transitive closure
  - The next slide shows a *manager* relation
  - Each step of the iterative process constructs an extended version of *empl* from its recursive definition.
  - The final result is called the *fixed point* of the recursive view definition.
- Recursive views are required to be *monotonic*. That is, if we add tuples to *manger* the view contains all of the tuples it contained before, plus possibly more





# Example of Fixed-Point Computation

| <i>employee_name</i> | <i>manager_name</i> |
|----------------------|---------------------|
| Alon                 | Barinsky            |
| Barinsky             | Estovar             |
| Corbin               | Duarte              |
| Duarte               | Jones               |
| Estovar              | Jones               |
| Jones                | Klinger             |
| Rensal               | Klinger             |

| <i>Iteration number</i> | <i>Tuples in empl</i>                             |
|-------------------------|---------------------------------------------------|
| 0                       |                                                   |
| 1                       | (Duarte), (Estovar)                               |
| 2                       | (Duarte), (Estovar), (Barinsky), (Corbin)         |
| 3                       | (Duarte), (Estovar), (Barinsky), (Corbin), (Alon) |
| 4                       | (Duarte), (Estovar), (Barinsky), (Corbin), (Alon) |





# Advanced SQL Features\*\*

- Create a table with the same schema as an existing table:  
`create table temp_account like account`
- SQL:2003 allows subqueries to occur *anywhere* a value is required provided the subquery returns only one value. This applies to updates as well
- SQL:2003 allows subqueries in the **from** clause to access attributes of other relations in the **from** clause using the **lateral** construct:

```
select C.customer_name, num_accounts
 from customer C,
 lateral (select count(*)
 from account A
 where A.customer_name = C.customer_name)
 as this_customer (num_accounts)
```





# Advanced SQL Features (cont'd)

- Merge construct allows batch processing of updates.
- Example: relation *funds\_received* (*account\_number*, *amount*) has batch of deposits to be added to the proper account in the *account* relation

```
merge into account as A
using (select *
 from funds_received as F)
on (A.account_number = F.account_number)
when matched then
 update set balance = balance + F.amount
```





# TRIGGERS

- CREATE OR REPLACE TRIGGER t  
BEFORE

**INSERT OR UPDATE OF salary, department\_id OR DELETE ON employees**  
**BEGIN CASE**

```
WHEN INSERTING THEN DBMS_OUTPUT.PUT_LINE('Inserting');
WHEN UPDATING('salary') THEN DBMS_OUTPUT.PUT_LINE('Updating salary');
WHEN UPDATING('department_id') THEN DBMS_OUTPUT.PUT_LINE('Updating
 department ID');
WHEN DELETING THEN DBMS_OUTPUT.PUT_LINE('Deleting'); END CASE; END;
/
```







creates a log table and a trigger that inserts a row in the log table after any **UPDATE** statement affects the **SALARY** column of the **EMPLOYEES** table, and then updates **EMPLOYEES.SALARY** and shows the log table.

- CREATE TABLE Emp\_log ( Emp\_id NUMBER, Log\_date DATE, New\_salary NUMBER, Action VARCHAR2(20));
  - Create trigger that inserts row in log table after EMPLOYEES.SALARY is update
- CREATE OR REPLACE TRIGGER log\_salary\_increase  
AFTER UPDATE OF salary ON employees  
FOR EACH ROW BEGIN INSERT INTO Emp\_log (Emp\_id, Log\_date, New\_salary, Action) VALUES (:NEW.employee\_id, SYSDATE, :NEW.salary, 'New Salary');  
END;





- Update EMPLOYEES.SALARY:

UPDATE employees SET salary = salary + 1000.0 WHERE  
Department\_id = 20; Result:

- 2 rows updated. Show log table:
- SELECT \* FROM Emp\_log;
- Result:

| EMP_ID | LOG_DATE  | NEW_SALARY | ACTION     |
|--------|-----------|------------|------------|
| 201    | 28-APR-10 | 15049.13   | New Salary |
| 202    | 28-APR-10 | 6945.75    | New Salary |

2 rows selected.





# End of Chapter

**Database System Concepts, 5th Ed.**

©Silberschatz, Korth and Sudarshan

See [www.db-book.com](http://www.db-book.com) for conditions on re-use

