

<b>Course Name:</b>	<b>Data Analysis Laboratory (216H03L501 )</b>	<b>Semester:</b>	<b>V</b>
<b>Date of Performance:</b>	<b>28 / 07 / 2025</b>	<b>DIV/ Batch No:</b>	<b>D2</b>
<b>Student Name:</b>	<b>Shreyans Tatiya</b>	<b>Roll No:</b>	<b>16010123325</b>

### Experiment No: 2

**Title:** Implement data pre-processing using python on real world dataset

**Objectives of the Experiment:**

To understand the need for data preprocessing before analysis or model building

**COs to be achieved:**

CO1: Understand basic concepts of data analytics to solve real-world problems

**Books/ Journals/ Websites referred:**

1. <https://pandas.pydata.org/>
2. <https://numpy.org/>
3. This experiment

**Theory:**

Pandas is a powerful open-source Python library used for data manipulation and analysis. It provides data structures like DataFrame and Series which allow efficient handling of structured data. With pandas, operations like filtering, grouping, merging, reshaping, and cleaning data become intuitive and fast.

**Problem statement/ Tasks**

**Program:**

```
import pandas as pd
```

```
import numpy as np
```

```
# Sample data
```

```

data = {
    'name': ['Alice', 'Bob', 'Charlie', 'Dave', 'Eve'],
    'age': [25, np.nan, 30, 22, 35],
    'gender': ['F', 'M', 'M', 'M', 'F'],
    'income': [50000, 60000, 75000, np.nan, 80000]
}

df = pd.DataFrame(data)

# Display the original data

print("Original DataFrame:")

print(df)

# User-defined function for discretization

def discretize_age(age):
    if age < 30:
        return 'Young'
    elif age >= 30 and age < 40:
        return 'Middle-aged'
    else:
        return 'Old'

# To view result

df['age_group'] = df['age'].apply(discretize_age)

# Handling missing values (NaN)

# Fill missing values in 'age' with the mean age

```

```

mean_age = df['age'].mean()

df['age'].fillna(mean_age, inplace=True)

# Apply discretization function to 'age' column

df['age_category'] = df['age'].apply(discretize_age)

# Drop rows with missing values in any column

df.dropna(inplace=True)

# Convert categorical variables (gender) to numerical

df['gender'] = df['gender'].map({'F': 0, 'M': 1})

# find employee with maximum salary

df.loc[df['income'].idxmax(), 'name']

#find youngest employee

df.loc[df['age'].idxmax(), 'name']

# Data normalization Min -Max

# Normalize 'income' column to range [0, 1]

min_income = df['income'].min()

max_income = df['income'].max()

df['income_normalized'] = (df['income'] - min_income) / (max_income - min_income)

# Display cleaned, preprocessed, and discretized data

print("\nCleaned, Preprocessed, and Discretized DataFrame:")

print(df)

```

**Task: Download the real time data set and implement data preprocessing techniques on the real time data set**

**Source of the dataset (URL):**

**Platform used by the student:**

**Following points should be written by students**

Different steps in Data Preprocessing:

- **Finding missing, null values**
- **Replacing missing, null values with statistical parameters**
- **Encoding categorical data if needed (Write user defined function)**
- **Normalization (Write user defined function)**
- **Discretization (Write user defined function)**

**Code :**

<https://colab.research.google.com/drive/146p8irTZTIqJb8HW3xQjYKUTnaHUviXp?usp=sharing>

**Output:**



The screenshot shows a Jupyter Notebook interface with the following details:

- Header:** Expt2.ipynb, File, Edit, View, Insert, Runtime, Tools, Help, Share, Gemini.
- Toolbar:** Reconnect, Run all.
- Code Cells:**
  - [ ] import pandas as pd
  - [ ] df = pd.read\_csv('housing - housing.csv')
  - [ ] df.head()
- Data Preview:** A table showing the first 5 rows of the 'housing' dataset.

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	median_house_value	ocean_proximity
0	-122.23	37.88	41.0	880	129.0	322.0	126.0	8.3252	452600	NEAR BAY
1	-122.22	37.86	21.0	7099	1106.0	2401.0	1138.0	8.3014	358500	NEAR BAY
2	-122.24	37.85	52.0	1467	190.0	496.0	177.0	7.2574	352100	NEAR BAY
3	-122.25	37.85	52.0	1274	235.0	558.0	219.0	5.6431	341300	NEAR BAY
4	-122.25	37.85	52.0	1627	280.0	565.0	259.0	3.8462	342200	NEAR BAY
- Bottom Navigation:** Next steps: Generate code with df, View recommended plots, New interactive sheet.

The screenshot shows a Jupyter Notebook interface with the following details:

- Header:** Expt2.ipynb, File, Edit, View, Insert, Runtime, Tools, Help, Share, Gemini.
- Toolbar:** Commands, Code, Text, Run all.
- Output Area:** Shows the result of the command `[ 1] print(df.isnull().sum())`, which prints a summary of missing values across various columns.
- Code Area:** Shows the following code snippets:
  - `[ 1] df['total_bedrooms'] = df['total_bedrooms'].fillna(df['total_bedrooms'].mean())`
  - `[ 1] df['housing_median_age'] = df['housing_median_age'].fillna(df['housing_median_age'].mean())`
  - `[ 1] df['population'] = df['population'].fillna(df['population'].mean())`
  - `[ 1] df['households'] = df['households'].fillna(df['households'].mean())`
  - `[ 2] # Mode produces a series type, we need to use [0] to get only the value, else it'll give us dtype and more areas we don't require`
  - `[ 2] df['ocean_proximity'] = df['ocean_proximity'].fillna(df['ocean_proximity'].mode()[0])`



```
Expt2.ipynb ⌂ ⌂
File Edit View Insert Runtime Tools Help
Q Commands + Code + Text > Run all ▾

[ ] print(df)

    longitude latitude housing_median_age total_rooms total_bedrooms \
0 -122.23 37.88 41.0 888 129.0
1 -122.22 37.86 21.0 7999 1106.0
2 -122.24 37.85 52.0 1467 190.0
3 -122.25 37.85 52.0 1274 235.0
4 -122.25 37.85 52.0 1627 280.0
...
28635 -121.99 36.48 25.0 1665 374.0
28636 -121.21 39.49 18.0 697 150.0
28637 -121.22 39.43 17.0 2254 485.0
28638 -121.32 39.43 18.0 1868 489.0
28639 -121.24 39.37 16.0 2785 616.0

population households median_income median_house_value \
0 322.0 126.0 8.3252 452600
1 2401.0 1138.0 8.3814 358500
2 496.0 177.0 7.2574 352100
3 558.0 219.0 5.6431 341300
4 565.0 259.0 3.8462 342200
...
28635 945.0 339.0 1.5563 78100
28636 356.0 114.0 2.3568 77100
28637 1807.0 433.0 1.7080 92300
28638 741.0 349.0 1.8672 84700
28639 1387.0 530.0 2.3886 89400

ocean_proximity
0 NEAR BAY
1 NEAR BAY
2 NEAR BAY
3 NEAR BAY
4 NEAR BAY
...
28635 INLAND
28636 INLAND
28637 INLAND
28638 INLAND
28639 INLAND

[20640 rows x 10 columns]

[ ] print(df.isnull().sum())

longitude      0
latitude       0
housing_median_age 0
total_rooms     0
total_bedrooms 0
population      0
households      0
median_income    0
median_house_value 0
ocean_proximity 0
dtype: int64
```

```
Expt2.ipynb ⌂ ⌂
File Edit View Insert Runtime Tools Help
Q Commands + Code + Text > Run all ▾

[ ] df['ocean_proximity_encoded'] = df['ocean_proximity'].astype('category').cat.codes

[ ] def discretize_income(income):
    if income < 2.5:
        return 'Low'
    elif income < 4.5:
        return 'Medium'
    else:
        return 'High'

df['income_category'] = df['median_income'].apply(discretize_income) # adding a column called 'income category' which helps us to convert continuous numerical data into categorical buckets

[ ] def normalize_column(col):
    return (col - col.min()) / (col.max() - col.min())

df['normalized_income'] = normalize_column(df['median_income']) # Scales continuous numerical data to a standard range - usually [0, 1].
```



The screenshot shows a Jupyter Notebook interface with the following content:

```
[ ] max_price_index = df['median_house_value'].idxmax() # returns the max value row from this column "median-house-value"
print("House with Maximum Value:")
print(df.loc[max_price_index])
```

Output:

```
House with Maximum Value:
longitude    -122.27
latitude     37.85
housing_median_age 52.0
total_rooms   249
total_bedrooms 78.0
population    395.0
households    88.0
median_income  1.2434
median_house_value 580001
ocean_proximity <1H OCEAN
ocean_proximity_encoded 0
income_category Low
normalized_income 0.051275
Name: 89, dtype: object
```

```
df.head()
```

Output:

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	median_house_value	ocean_proximity	ocean_proximity_encoded	income_category	normalized_income
0	-122.23	37.88	41.0	880	129.0	322.0	126.0	8.3252	452600	NEAR BAY	3	High	0.539668
1	-122.22	37.86	21.0	7099	1106.0	2401.0	1138.0	8.3014	358500	NEAR BAY	3	High	0.538627
2	-122.24	37.85	52.0	1467	190.0	496.0	177.0	7.2574	352100	NEAR BAY	3	High	0.466028
3	-122.25	37.85	52.0	1274	235.0	558.0	219.0	5.6431	341300	NEAR BAY	3	High	0.354699
4	-122.25	37.85	52.0	1627	280.0	565.0	259.0	3.8462	342200	NEAR BAY	3	Medium	0.230776

### Post Lab Subjective/Objective type Questions:

**Q.1 What are some common challenges encountered during data cleaning? How did you handle missing values in the provided dataset?**

**Ans:**

Common challenges in data cleaning include:

- Missing values
- Inconsistent data types
- Outliers
- Duplicate entries

In the provided dataset, we handled missing values by:

Using the mode (most frequent value) to fill missing entries in the ocean\_proximity column:  
df['ocean\_proximity'] = df['ocean\_proximity'].fillna(df['ocean\_proximity'].mode()[0])

**Q.2 Explain the importance of data normalization in the context of machine learning models. How does normalizing benefit the analysis?**

**Ans:**

- Normalization scales numerical features to a common range (typically between 0 and 1).
- Without normalization, features with larger values can dominate those with smaller values, leading to biased results.
- In this lab, normalization helped us compare median\_income values more effectively:  

$$\text{df['normalized_income']} = (\text{df['median_income']} - \text{df['median_income'].min()}) / (\text{df['median_income'].max()} - \text{df['median_income'].min()})$$

**Q.3 Discuss why it's essential to convert categorical variables like 'gender' into numerical representations.**

**Ans:** Categorical variables (e.g., gender, ocean\_proximity) contain text labels that machine learning models **cannot process directly**. Converting them to numerical form allows:

- Better compatibility with models
- Logical comparisons and distance calculations

For example, we used **label encoding** to convert the ocean\_proximity column:

```
df['ocean_proximity_encoded'] = df['ocean_proximity'].astype('category').cat.codes
```

This makes the data model-friendly while preserving category information.

### **Conclusion:**

In this lab, I learned how essential data cleaning and preprocessing steps are for preparing real-world datasets. I handled missing values, encoded categorical features, normalized numerical data, and extracted useful insights.