

Introduction to Reinforcement Learning (RL)

[Vikky Masih,](#)


Research Scholar,

Mehta Family School of Data Science & Artificial Intelligence,

IIT Guwahati


Machine Learning (ML)

- ML is fundamental concept of AI
 - Learn to improve performance via experience




- Unsupervised Learning:
 - Clustering, Anomaly Detection, PCA, etc.
 - Find patterns in input data
- Supervised Learning:
 - Classification and Regression
 - Labelled data for training
- Reinforcement Learning:
 - Decision making under uncertainty
 - Learn to improve performance via interacting with environment

What do you want the machine learning system to do?




Credit: Thomas Malone, MIT Sloan | Design: Laura Wentzel

Machine Learning (ML)




- Unsupervised Learning:
 - Clustering, Anomaly Detection, PCA, etc.
 - Find patterns in input data
- Unsupervised learning example:
 - Goal: Clustering, Outlier detection
 - Data: $[\vec{x}_1, \dots, \vec{x}_n]$

Machine Learning (ML)




Classification



Regression


- Supervised Learning:
 - Classification and Regression
 - Labelled data for training
- Regression:
 - Goal: $f(\vec{x}) = \vec{y}$,
 - Data: $[(\vec{x}_i, \vec{y}_i), \dots, (\vec{x}_n, \vec{y}_n)]$
- Classification:
 - Goal: $\operatorname{argmax} P(\text{class}|\vec{x}) = C$
 - Data: $[(\vec{x}_i, C_i), \dots, (\vec{x}_n, C_n)]$

Machine Learning (ML)



- Reinforcement Learning:
 - Stochastic optimal control
 - Learn to improve performance via interacting with the environment
- RL example:
 - Goal:
 - Maximize cumulative reward
 $\text{Maximize } \sum_{i=1}^{\infty} \text{Reward}(\text{State}_i, \text{Act}_i)$
 - Data:
 $\text{Reward}_{i+1}, \text{State}_{i+1} = \text{Interact}(\text{State}_i, \text{Action}_i)$


TD-Gammon – Tesauro ~1995




- Net with 80 hidden units, initialize to random weights
- Select move based on network estimate & shallow search
- Learn by playing against itself
- 1.5 million games of training
 - competitive with world class players

Atari 2600 games

State: Raw Pixels
Actions: Valid Moves
Reward: Game Score



Same model/parameters for
~50 games



Robotics and Locomotion

State:

Joint States/Velocities

Accelerometer/Gyroscope

Terrain

Actions: Apply Torque to Joints

Reward: Velocity – { stuff }







Figure 5: Time-lapse images of a representative *Quadruped* policy traversing gaps (left); and navigating obstacles (right)






https://youtu.be/hx_bgoTF7bs




Alpha Go

- Learning how to beat humans at 'hard' games (search space too big)
- Far surpasses (Human) Supervised learning
- Algorithm learned to outplay humans at chess in 24 hours

State: Board State
Actions: Valid Moves
Reward: Win or Lose






Agent and environment interact at discrete time steps: $t = 0, 1, 2, 3, \dots$

Agent observes state at step t : $S_t \in \mathcal{S}$

produces action at step t : $A_t \in \mathcal{A}(S_t)$

gets resulting reward: $R_{t+1} \in \mathcal{R} \subset \mathbb{R}$


and resulting next state: $S_{t+1} \in \mathcal{S}^+$



Reinforcement Learning

- RL problems are generally posed as **Markov Decision Process** (MDP)
 - RL is used for MDPs where the transition prob. or reward prob. are unknown.
- MDP: **Discrete-Time Stochastic Control Process**
 - Markovian Property:
 - Next reward and state **does not** depend on **history**.
 - Next reward and state **depend** only on **current state and action**.
 - It's a 4-tuple $(S, A_s, P_a(s, s'), R_a(s, s'))$
 - $S \rightarrow$ State Space $\rightarrow$ Set of states
 - $A_s \rightarrow$ Action Space available at state **s** $\rightarrow$ Set of possible actions
 - $P_a(s, s') = \mathbb{P}(s'|s, a) \rightarrow$ Probability of transitioning from **s** to **s'** after taking action **a**
 - $R_a(s, s') = \mathbb{E}(r'|s, a) \rightarrow$ Expected Reward obtained **after** transitioning from **s** to **s'** after taking action **a**

Example MDP




- 3 states (green circles)
 - $\{S_0, S_1, S_2\}$
- 2 actions (orange circles)
 - $\{a_0, a_1\}$
- 2 rewards (orange arrows)
 - $R_{a_1}(S_2, S_0) = -1$
 - $R_{a_0}(S_1, S_0) = +5$
- Example transition probabilities:
 - $P_{a_0}(S_0, S_2) = 0.5$
 - $P_{a_0}(S_0, S_0) = 0.5$

Reinforcement Learning

- Policy (π): Mapping from state to action
 - Deterministic: $a_t = \pi(s_t)$, or,
 - Stochastic: $a_t = \operatorname{argmax}_a \pi(a|s_t)$
- Objective of RL:
 - Find a policy that maximizes long term cumulative reward.
 - maximize $\sum_{t=0}^T \gamma^t R_{a_t}(s_t, s_{t+1})$, where, $\gamma \in [0,1]$ (Discount factor)
- How to make a decision?
 - Rank State or (State, Action) based on some value derived from experience
 - Value functions measure the goodness of a particular state or state/action pair for a given Policy

Robot in a room



- States:
 - Location $\in \{(1,1), \dots, (3,4)\}$
- Actions:
 - UP, DOWN, LEFT, RIGHT
- Terminate at (1,4) or (2,4)
- Note:
 - Transitions and rewards are deterministic.
- Reward +1 at (1,4), -1 at (2,4)
- Reward -0.1 for each step

Robot in a room: State Value Function

?	?	?	+1
?		?	-1
?	?	?	?

Reward -0.1 for each step

- State Value Function:
 - $V(s)$
 - Maximum expected reward accumulated when starting from a given state
- Bellman equation (Optimal):
 - $V(s) = \max_a \left(\mathbb{E}(R_a(s, s') + \gamma V(s')) \right)$
 - $\gamma = 1, s = s_t, s' = s_{t+1}$

Robot in a room: State Value Function

0	0	0	+1
0		0	-1
0	0	0	0

Reward -0.1 for each step

- State Value Function:
 - $V(s)$
 - Maximum expected reward accumulated when starting from a given state
- Bellman equation (Optimal):
 - $V(s) = \max_a \left(\mathbb{E}(R_a(s, s') + \gamma V(s')) \right)$
 - $\gamma = 1, s = s_t, s' = s_{t+1}$
- Value Iteration
 - Initializing
 - $V_{k+1}(s) = \max_a \left(\mathbb{E}(R_a(s, s') + \gamma V_k(s')) \right)$

Robot in a room: State Value Function

-0.1	-0.1	0.9	+1
-0.1		-0.1	-1
-0.1	-0.1	-0.1	-0.1

Reward -0.1 for each step

- State Value Function:
 - $V(s)$
 - Maximum expected reward accumulated when starting from a given state
- Bellman equation (Optimal):
 - $V(s) = \max_a \left(\mathbb{E}(R_a(s, s') + \gamma V(s')) \right)$
 - $\gamma = 1, s = s_t, s' = s_{t+1}$
- Using Bellman equation iteratively
 - $V_{k+1}(s) = \max_a \left(\mathbb{E}(R_a(s, s') + \gamma V_k(s')) \right)$
 - First Iteration

Robot in a room: State Value Function

-0.1	0.8	0.9	+1
-0.1		0.8	-1
-0.1	-0.1	-0.1	-0.1

Reward -0.1 for each step

- State Value Function:
 - $V(s)$
 - Maximum expected reward accumulated when starting from a given state
- Bellman equation (Optimal):
 - $V(s) = \max_a \left(\mathbb{E}(R_a(s, s') + \gamma V(s')) \right)$
 - $\gamma = 1, s = s_t, s' = s_{t+1}$
- Using Bellman equation iteratively
 - $V_{k+1}(s) = \max_a \left(\mathbb{E}(R_a(s, s') + \gamma V_k(s')) \right)$
 - Second Iteration

Robot in a room: State Value Function

0.7	0.8	0.9	+1
-0.1		0.8	-1
-0.1	-0.1	0.7	-0.1

Reward -0.1 for each step

- State Value Function:
 - $V(s)$
 - Maximum expected reward accumulated when starting from a given state
- Bellman equation (Optimal):
 - $V(s) = \max_a \left(\mathbb{E}(R_a(s, s') + \gamma V(s')) \right)$
 - $\gamma = 1, s = s_t, s' = s_{t+1}$
- Using Bellman equation iteratively
 - $V_{k+1}(s) = \max_a \left(\mathbb{E}(R_a(s, s') + \gamma V_k(s')) \right)$
 - Third Iteration

Robot in a room: State Value Function

0.7	0.8	0.9	+1
0.6		0.8	-1
-0.1	0.6	0.7	0.6

Reward -0.1 for each step

- State Value Function:
 - $V(s)$
 - Maximum expected reward accumulated when starting from a given state
- Bellman equation (Optimal):
 - $V(s) = \max_a \left(\mathbb{E}(R_a(s, s') + \gamma V(s')) \right)$
 - $\gamma = 1, s = s_t, s' = s_{t+1}$
- Using Bellman equation iteratively
 - $V_{k+1}(s) = \max_a \left(\mathbb{E}(R_a(s, s') + \gamma V_k(s')) \right)$
 - Fourth Iteration

Robot in a room: State Value Function

0.7	0.8	0.9	+1
0.6		0.8	-1
0.5	0.6	0.7	0.6

Reward -0.1 for each step

- State Value Function:
 - $V(s)$
 - Maximum expected reward accumulated when starting from a given state
- Bellman equation (Optimal):
 - $V(s) = \max_a \left(\mathbb{E}(R_a(s, s') + \gamma V(s')) \right)$
 - $\gamma = 1, s = s_t, s' = s_{t+1}$
- Using Bellman equation iteratively
 - $V_{k+1}(s) = \max_a \left(\mathbb{E}(R_a(s, s') + \gamma V_k(s')) \right)$
 - Fifth Iteration


Robot in a room: State Value Function

0.7	0.8	0.9	+1
0.6		0.8	-1
0.5	0.6	0.7	0.6

Reward -0.1 for each step

- State Value Function:
 - $V(s)$
 - Maximum expected reward accumulated when starting from a given state
- Bellman equation (Optimal):
 - $V(s) = \max_a \left(\mathbb{E}(R_a(s, s') + \gamma V(s')) \right)$
 - $\gamma = 1, s = s_t, s' = s_{t+1}$
- Using Bellman equation iteratively
 - $V_{k+1}(s) = \max_a \left(\mathbb{E}(R_a(s, s') + \gamma V_k(s')) \right)$
 - Converged

Robot in a room: State Value Function



Reward -0.1 for each step

- State Value Function:
 - $V(s)$
 - Maximum expected reward accumulated when starting from a given state
- Bellman equation (Optimal):
 - $V(s) = \max_a \left(\mathbb{E}(R_a(s, s') + \gamma V(s')) \right)$
 - $\gamma = 1, s = s_t, s' = s_{t+1}$
- Policy:
 - $\pi(s) = \operatorname{argmax}_a \sum P_a(s, s') V(s')$

Robot in a room:

?	?	?	+1
?		?	-1
?	?	?	?

Reward -0.1 for each step

What if the robot is not functioning properly?

- State transitions are stochastic
 - An action may not lead to intended state
- Rewards/Costs are stochastic

Robot in a room: State-Action Value Function

→ ? ↓ ?	← ? → ?	← ? → ? ↓ ?	+1
↑ ? ↓ ?		↑ ? → ? ↓ ?	-1
↑ ? → ?	← ? → ?	↑ ? ← ? → ?	↑ ? ← ?

Reward -0.1 for each step

- State-Action Value Function:

- $Q(s, a)$
- Maximum expected reward accumulated when starting from a given state and choosing a given action.

- $V(s) = \max_a Q(s, a)$

- Bellman equation (Optimal):

$$Q(s, a) = \mathbb{E} \left(R_a(s, s') + \gamma \max_{a'} Q(s', a') \right)$$

$\gamma = 1$, $s = s_t$, $s' = s_{t+1}$, $a' = a_{t+1}$

Robot in a room: State-Action Value Function

→ 0 ↓ 0	← 0 → 0	← 0 → 0 ↓ 0	+1
↑ 0 ↓ 0		↑ 0 → 0 ↓ 0	-1
↑ 0 → 0	← 0 → 0	↑ 0 ← 0 → 0	↑ 0 ← 0

Reward -0.1 for each step

- State-Action Value Function:

- $Q(s, a)$
- Maximum expected reward accumulated when starting from a given state and choosing a given action.

- $V(s) = \max_a Q(s, a)$

- Bellman equation (Optimal):

$$Q(s, a) = \mathbb{E} \left(R_a(s, s') + \gamma \max_{a'} Q(s', a') \right)$$

$$\gamma = 1, s = s_t, s' = s_{t+1}, a' = a_{t+1}$$

- Value Iteration: Initialization

Robot in a room: State-Action Value Function

→ -0.1 ↓ -0.1	← -0.1 → -0.1	← -0.1 → 0.9 ↓ -0.1	+1
↑ -0.1 ↓ -0.1		↑ -0.1 → -1.1 ↓ -0.1	-1
↑ -0.1 → -0.1	← -0.1 → -0.1	↑ -0.1 ← -0.1 → -0.1	↑ -1.1 ← -0.1

Reward -0.1 for each step

- State-Action Value Function:

- $Q(s, a)$
- Maximum expected reward accumulated when starting from a given state and choosing a given action.

- $V(s) = \max_a Q(s, a)$

- Bellman equation (Optimal):

$$Q(s, a) = \mathbb{E} \left(R_a(s, s') + \gamma \max_{a'} Q(s', a') \right)$$

$$\gamma = 1, s = s_t, s' = s_{t+1}, a' = a_{t+1}$$

- Value Iteration: First Iteration

Robot in a room: State-Action Value Function

→ -0.2 ↓ -0.2	← -0.2 → 0.8	← -0.2 → 0.9 ↓ -0.2	+1
↑ -0.2 ↓ -0.2		↑ 0.8 → -1.1 ↓ -0.2	-1
↑ -0.2 → -0.2	← -0.2 → -0.2	↑ -0.2 ← -0.2 → -0.2	↑ -1.1 ← -0.2

Reward -0.1 for each step

- State-Action Value Function:

- $Q(s, a)$
- Maximum expected reward accumulated when starting from a given state and choosing a given action.

- $V(s) = \max_a Q(s, a)$

- Bellman equation (Optimal):

$$Q(s, a) = \mathbb{E} \left(R_a(s, s') + \gamma \max_{a'} Q(s', a') \right)$$

$$\gamma = 1, s = s_t, s' = s_{t+1}, a' = a_{t+1}$$

- Value Iteration: Second Iteration

Robot in a room: State-Action Value Function

→ 0.7 ↓ -0.3	← -0.3 → 0.8	← 0.7 → 0.9 ↓ 0.7	+1
↑ -0.3 ↓ -0.3		↑ 0.8 → -1.1 ↓ -0.3	-1
↑ -0.3 → -0.3	← -0.3 → -0.3	↑ 0.7 ← -0.3 → -0.3	↑ -1.1 ← -0.3

Reward -0.1 for each step

- State-Action Value Function:

- $Q(s, a)$
- Maximum expected reward accumulated when starting from a given state and choosing a given action.

- $V(s) = \max_a Q(s, a)$

- Bellman equation (Optimal):

$$Q(s, a) = \mathbb{E} \left(R_a(s, s') + \gamma \max_{a'} Q(s', a') \right)$$

$$\gamma = 1, s = s_t, s' = s_{t+1}, a' = a_{t+1}$$

- Value Iteration: Third Iteration

Robot in a room: State-Action Value Function

→ 0.7 ↓ 0.5	← 0.6 → 0.8	← 0.7 → 0.9 ↓ 0.7	+1
↑ 0.6 ↓ 0.4		↑ 0.8 → -1.1 ↓ 0.6	-1
↑ 0.5 → 0.5	← 0.4 → 0.6	↑ 0.7 ← 0.5 → 0.5	↑ -1.1 ← 0.6

Reward -0.1 for each step

- State-Action Value Function:

- $Q(s, a)$
- Maximum expected reward accumulated when starting from a given state and choosing a given action.


- $V(s) = \max_a Q(s, a)$

- Bellman equation (Optimal):

$$Q(s, a) = \mathbb{E} \left(R_a(s, s') + \gamma \max_{a'} Q(s', a') \right)$$

$$\gamma = 1, s = s_t, s' = s_{t+1}, a' = a_{t+1}$$

Robot in a room: State-Action Value Function





Reward -0.1 for each step

- State-Action Value Function:
 - $Q(s, a)$
 - Maximum expected reward accumulated when starting from a given state and choosing a given action.
 - $V(s) = \max_a Q(s, a)$
- Policy:
 - $\pi(s) = \operatorname{argmax}_a Q(s, a)$

Policy Gradient Method

- Policy Gradient:
 - learn policy directly $\pi(a|s)$



Policy Gradient Method

- Policy Gradient:
 - learn policy directly $\pi(a|s)$


- REINFORCE Algorithm

Initialize θ arbitrarily

for each episode $\{s_0, a_0, r_0, \dots, s_T, a_T, r_T\} \sim \pi(\cdot, \cdot; \theta)$

do $\left\{ \begin{array}{l} \textbf{for } t \leftarrow 0 \textbf{ to } T \\ \textbf{do} \left\{ \begin{array}{l} G \leftarrow \sum_{k=t}^T \gamma^{k-t} \cdot r_k \\ \theta \leftarrow \theta + \alpha \cdot \gamma^t \cdot \nabla_{\theta} \log \pi(s_t, a_t; \theta) \cdot G \end{array} \right. \end{array} \right.$

return (θ)



References

- “What is reinforcement learning?,” What Is Reinforcement Learning? - MATLAB & Simulink, <https://www.mathworks.com/discovery/reinforcement-learning.html> (accessed Jul. 10, 2023).
- S. Brown, “Machine Learning, explained,” MIT Sloan, <https://mitsloan.mit.edu/ideas-made-to-matter/machine-learning-explained> (accessed Jul. 10, 2023).
- “Cluster analysis,” Wikipedia, https://en.wikipedia.org/wiki/Cluster_analysis (accessed Jul. 10, 2023).
- “Regression vs classification in Machine Learning - Javatpoint,” www.javatpoint.com, <https://www.javatpoint.com/regression-vs-classification-in-machine-learning> (accessed Jul. 10, 2023).
- “Reinforcement learning,” Wikipedia, https://en.wikipedia.org/wiki/Reinforcement_learning (accessed Jul. 10, 2023).
- G. Hulten, “CSEP546: Machine learning,” University of Washington, <https://courses.cs.washington.edu/courses/csep546/18au/> (accessed Jul. 10, 2023).

References

- K. Fragkiadaki and T. Mitchell, “CMU 10703: Deep RL and Control,” Carnegie Mellon University, <https://www.andrew.cmu.edu/course/10-703/> (accessed Jul. 10, 2023).
- “Markov decision process,” Wikipedia, https://en.wikipedia.org/wiki/Markov_decision_process (accessed Jul. 10, 2023).
- P. Bodik, “Practical machine learning lecture: Reinforcement learning,” University of California, Berkeley, <http://people.eecs.berkeley.edu/~jordan/courses/294-fall09/lectures/reinforcement/> (accessed Jul. 10, 2023).
- A. Karpathy et al., “Deep RL Bootcamp,” Deep RL Bootcamp - Berkeley CA, <https://sites.google.com/view/deep-rl-bootcamp/lectures> (accessed Jul. 10, 2023).
- R. S. Sutton and A. Barto, Reinforcement Learning: An Introduction. Cambridge, Massachusetts ; London, England: The MIT Press, 2020.

Thankyou

Appendices

A1. Optimal State-Value Function

- Value function for arbitrary π

$$\begin{aligned}v_{\pi}(s) &\doteq \mathbb{E}_{\pi}[G_t | S_t = s] \\&= \mathbb{E}_{\pi}[R_{t+1} + \gamma G_{t+1} | S_t = s] \\&= \sum_a \pi(a|s) \sum_{s', r} p(s', r | s, a) [r + \gamma v_{\pi}(s')]\end{aligned}$$

- Optimal value function

$$\begin{aligned}v_*(s) &\doteq \max_{\pi} v_{\pi}(s) \\&= \max_a \mathbb{E}[R_{t+1} + \gamma v_*(S_{t+1}) | S_t = s, A_t = a] \\&= \max_a \sum_{s', r} p(s', r | s, a) [r + \gamma v_*(s')]\end{aligned}$$

- Return

$$\begin{aligned}G_t &\doteq R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots \\&= \sum_{k=0}^{\infty} \gamma^k R_{t+k+1} \\&= R_{t+1} + \gamma G_{t+1}\end{aligned}$$

A2. Optimal (State, Action)-Value Function

- Q function for arbitrary π

$$\begin{aligned} q_{\pi}(s, a) &\doteq \mathbb{E}_{\pi}[G_t | S_t = s, A_t = a] \\ &= \sum_{s', r} p(s', r | s, a) \left[r + \gamma \sum_{a'} \pi(a' | s') q_{\pi}(a', s') \right] \end{aligned}$$

- Optimal Q function

$$\begin{aligned} q_*(s, a) &\doteq \max_{\pi} q_{\pi}(s, a) \\ &= \mathbb{E}[R_{t+1} + \gamma \max_{a'} q_*(S_{t+1}, a') | S_t = s, A_t = a] \\ &= \sum_{s', r} p(s', r | s, a) \left[r + \gamma \max_{a'} q_*(s', a') \right] \end{aligned}$$

- Return

$$\begin{aligned} G_t &\doteq R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots \\ &= \sum_{k=0}^{\infty} \gamma^k R_{t+k+1} \\ &= R_{t+1} + \gamma G_{t+1} \end{aligned}$$

A3. Value iteration

Params: θ - a small positive threshold determining the accuracy of the estimation

Initialize $V(s)$, for all $s \in \mathcal{S}^+$ arbitrarily, except $V(\text{terminal})$

$\Delta \leftarrow 0$

while $\Delta \geq \theta$ **do**

foreach $s \in S$ **do**

$v \leftarrow V(s)$

$V(s) \leftarrow \max_a \sum_{s', r} p(s', r | s, a) [r + \gamma V(s')]$

$\Delta \leftarrow \max(\Delta, |v - V(s)|)$

end

end

output: Deterministic policy $\pi \approx \pi_*$ such that

$\pi(s) = \operatorname{argmax}_a \sum_{s', r} p(s', r | s, a) [r + \gamma V(s')]$

Reinforcement Learning

Reinforcement learning

- Reinforcement Learning is the agent must sense the environment, learns to behave (act) in a environment by performing actions (reinforcement) and seeing the results.
- Task
 - Learn how to behave successfully to achieve a goal while interacting with an external environment.
 - The goal of the agent is to **learn** an **action policy** that maximizes the total reward it will receive from any starting state.
- Examples
 - **Game playing:** player knows whether it win or lose, but not know how to move at each step


Applications

- A robot cleaning the room and recharging its battery
- Robot-soccer
- invest in shares
- Modeling the economy through rational agents
- Learning how to fly a helicopter
- Scheduling planes to their destinations

Reinforcement Learning Process

- RL contains two primary components:
 1. Agent (A) – RL algorithm that learns from trial and error
 2. Environment – World Space in which the agent moves (interact and take action)
- State (S) – Current situation returned by the environment
- Reward (R) – An immediate return from the environment to appraise the last action
- Policy (π) – Agent uses this approach to decide the next action based on the current state
- Value (V) – Expected long-term return with discount. Oppose to the short-term reward (R)
- Action-Value (Q) – Similar to value except it contains an additional parameter, the current action (A)

Figure shows RL is learning from interaction



RL Approaches


- Two approaches
 - *Model based approach RL:*
 - learn the model, and use it to derive the optimal policy.
e.g Adaptive dynamic learning(ADP) approach
 - *Model free approach RL:*
 - derive the optimal policy without learning the model.
e.g LMS and Temporal difference approach
- Passive learning
 - The agent simply watches the world during transition and tries to learn the utilities in various states
- Active learning
 - The agent not simply watches, but also acts on the environment

Reinforcement learning model

- Each percept(e) is enough to determine the State(the state is accessible)
- **Agent's task:** Find a optimal policy by mapping states of environment to actions of the agent, that maximize long-run measure of the reward (reinforcement)
- It can be modeled as Markov Decision Process (MDP) model.
 - Markov decision process (**MDP**) is a a mathematical framework for **modeling** decision making i.e mapping a solution in reinforcement learning.

MDP model

- MDP model $\langle S, T, A, R \rangle$



- S — set of states
- A — set of actions
- Transition Function:** $T(s, a, s') = P(s'|s, a)$ — the probability of transition from s to s' given action a
 $T(s, a) \rightarrow s'$
- Reward Function:** $r(s, a) \rightarrow r$ the expected reward for taking action a in state s

$$R(s, a) = \sum_{s'} P(s'|s, a) r(s, a, s')$$

$$R(s, a) = \sum_{s'} T(s, a, s') r(s, a, s')$$

Q - Learning

- **Q-Learning** is a value-based **reinforcement learning** algorithm uses Q-values (action values) to iteratively improve the behavior of the learning agent.
- Goal is to maximize the Q value to find the optimal action-selection policy.
- The Q table helps to find the best action for each state and maximize the expected reward.
- **Q-Values / Action-Values:** Q-values are defined for states and actions.
- $Q(s, a)$ denotes an estimation of the action a at the state s .
- This estimation of $Q(s, a)$ will be iteratively computed using the **TD-Update rule**.
- **Reward:** At every transition, the agent observes a reward for every action from the environment, and then transits to another state.
- **Episode:** If at any point of time the agent ends up in one of the terminating states i.e. there are no further transition possible is called **completion of an episode**.

Q-Learning

- Initially agent explore the environment and update the Q-Table. When the Q-Table is ready, the agent will start to exploit the environment and taking better actions.
- It is an off-policy control algorithm i.e. the updated policy is different from the behavior policy. It estimates the reward for future actions and appends a value to the new state without any greedy policy. Temporal difference(TD) learning is one of the model-free reinforcement learning techniques

Temporal Difference or TD-Update:

- Estimate the value of Q is applied at every time step of the agents interaction with the environment. $Q(S,A) \leftarrow Q(S,A) + \alpha (R + \gamma Q(S',A') - Q(S,A))$

Advantage:

- Converges to an optimal policy in both deterministic and nondeterministic MDPs.

Disadvantage:

- Suitable for small problems

Bellman Equation

In the context of Q-learning, Bellman's equation is utilized to determine the value of a specific state and evaluate its relative placement.

The optimal state is determined by the state with the highest value.

The Bellman's equation can be represented as

$$Q(\text{state}, \text{action}) = R(\text{state}, \text{action}) + \text{Gamma} * \max[Q(\text{next state}, \text{all actions})]$$

Where,

$Q(s,a)$ represents the expected reward for an action 'a' in state 's'.


$R(s,a)$ represents the reward earned when action a is carried out in state 's'.

α is the discount factor, which denotes the significance of future rewards.

$\max_a Q(s',a)$ represents the maximum Q-value for the next state s' and every possible action.



Understanding the Q – Learning

- Building Environment contains 5 rooms that are connected with doors.
- Each room is numbered from 0 to 4. The building outside is numbered as 5.
- Doors from room 1 and 4 leads to the building outside 5.
- **Problem:** Agent can place at any one of the rooms (0, 1, 2, 3, 4). Agent's goal is to reach the building outside (room 5).




Understanding the Q – Learning

- Represent the room in the graph.
- Room number is the state and door is the edge.





Understanding the Q – Learning

- Assign the Reward value to each door.
- The doors lead immediately to target is assigned an instant reward of 100.
- Other doors not directly connected to the target room have **zero** reward.
- For example, doors are two-way (0 leads to 4, and 4 leads back to 0), two edges are assigned to each room.
- Each edge contains an instant reward value



Understanding the Q – Learning

- Let consider agent starts from state s (Room) 2.
- Agent's movement from one state to another state is action a .
- Agent is traversing from state 2 to state 5 (Target).
 - Initial state = current state i.e. state 2
 - Transition State 2 $\rightarrow$ State 3
 - Transition State 3 $\rightarrow$ State (2, 1, 4)
 - Transition State 4 $\rightarrow$ State 5




Understanding the Q – Learning

Prepare rewards table R (matrix)

- -1 denotes the no edge between the states
- 0 represents the indirect edge to the target

$$R = \begin{bmatrix} -1 & -1 & -1 & -1 & 0 & -1 \\ -1 & -1 & -1 & 0 & -1 & 100 \\ -1 & -1 & -1 & 0 & -1 & -1 \\ -1 & 0 & 0 & -1 & 0 & -1 \\ 0 & -1 & -1 & 0 & -1 & 100 \\ -1 & 0 & -1 & -1 & 0 & 100 \end{bmatrix}$$



Understanding the Q – Learning: **Prepare matrix Q**

- Matrix Q is the memory of the agent in which learned information from experience is stored.
- Row denotes the current state of the agent
- Column denotes the possible actions leading to the next state

Compute Q matrix:

$$Q(\text{state}, \text{action}) = R(\text{state}, \text{action}) + \text{Gamma} * \max[Q(\text{next state}, \text{all actions})]$$

- Gamma is discounting factor for future rewards. Its range is 0 to 1. i.e. $0 < \text{Gamma} < 1$.
- Future rewards are less valuable than current rewards so they must be discounted.
- If Gamma is closer to 0, the agent will tend to consider only the immediate rewards.
- If Gamma is closer to 1, the agent will tend to consider only future rewards with higher edge weights.

Q – Learning Algorithm

- Set the gamma parameter
- Set environment rewards in matrix R
- Initialize matrix Q as Zero
 - Select random initial (source) state
 - Set initial state s = current state
 - Select one action a among all possible actions using exploratory policy
 - Take this possible action a , going to the *next state* s' .
 - Observe reward r
 - Get maximum Q value to go to next state based on all possible actions
- Compute:
 - $Q(\text{state}, \text{action}) = R(\text{state}, \text{action}) + \text{Gamma} * \max[Q(\text{next state}, \text{all actions})]$
- Repeat the above steps until reach the goal state i.e current state = goal state

Example: Q – Learning

Matrix Q :


- Set the Gamma value = 0.8
- Initialize the matrix Q to zero matrix

$$Q = \begin{matrix} & \begin{matrix} 0 & 1 & 2 & 3 & 4 & 5 \end{matrix} \\ \begin{matrix} 0 \\ 1 \\ 2 \\ 3 \\ 4 \\ 5 \end{matrix} & \begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix} \end{matrix}$$

$$R = \begin{matrix} & \begin{matrix} \text{Action} \\ 0 & 1 & 2 & 3 & 4 & 5 \end{matrix} \\ \begin{matrix} \text{State} \\ 0 \\ 1 \\ 2 \\ 3 \\ 4 \\ 5 \end{matrix} & \begin{bmatrix} -1 & -1 & -1 & -1 & 0 & -1 \\ -1 & -1 & -1 & 0 & -1 & 100 \\ -1 & -1 & -1 & 0 & -1 & -1 \\ -1 & 0 & 0 & -1 & 0 & -1 \\ 0 & -1 & -1 & 0 & -1 & 100 \\ -1 & 0 & -1 & -1 & 0 & 100 \end{bmatrix} \end{matrix}$$

Example: Q – Learning

- From state 1, agent can go to either state 3 or 5.
- Let's choose 5.
- From 5, Compute Max Q value to go to
- Next state based on all possible actions



$$Q(\text{state}, \text{action}) = R(\text{state}, \text{action}) + \text{Gamma} * \max[Q(\text{next state}, \text{all actions})]$$


- $Q(1,5) = R(1,5) + 0.8 * \max[Q(5,1), Q(5,4), Q(5,5)]$
 $= 100 + 0.8 * \max[0, 0, 0] = 100 + 0 = \mathbf{100}$

- Update the Matrix Q.

$$R = \begin{matrix} & \begin{matrix} 0 & 1 & 2 & 3 & 4 & 5 \end{matrix} \\ \begin{matrix} 0 \\ 1 \\ 2 \\ 3 \\ 4 \\ 5 \end{matrix} & \begin{bmatrix} -1 & -1 & -1 & -1 & 0 & -1 \\ -1 & -1 & -1 & 0 & -1 & 100 \\ -1 & -1 & -1 & 0 & -1 & 100 \\ -1 & 0 & 0 & -1 & 0 & -1 \\ 0 & -1 & -1 & 0 & -1 & 100 \\ -1 & 0 & -1 & -1 & 0 & 100 \end{bmatrix} \end{matrix}$$

$$Q = \begin{matrix} & \begin{matrix} \text{State } 0 & 1 & 2 & 3 & 4 & 5 \end{matrix} \\ \begin{matrix} 0 \\ 1 \\ 2 \\ 3 \\ 4 \\ 5 \end{matrix} & \begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & \mathbf{100} \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix} \end{matrix}$$

- For next episode, choose next state 3 randomly that becomes current state.
- State 3 contains 3 choices i.e. state 1, 2 or 4.
- Let's choose state 1.
- Compute max Q value to go to next state based on all possible actions.



$$Q(\text{state}, \text{action}) = R(\text{state}, \text{action}) + \text{Gamma} * \max[Q(\text{next state}, \text{all actions})]$$

- $Q(3,1) = R(3,1) + 0.8 * \max[Q(1,3), Q(1,5)]$
 $= 0 + 0.8 * \max[0, 100] = 0 + 80 = \mathbf{80}$

- Update the Matrix Q.

							Action						
							State	0	1	2	3	4	5
$R =$	0	-1	-1	-1	0	-1	0	0	0	0	0	0	0
	1	-1	-1	-1	0	-1	1	0	0	0	0	0	100
	2	-1	-1	-1	0	-1	2	0	0	0	0	0	0
	3	-1	0	0	-1	0	3	0	80	0	0	0	0
	4	0	-1	-1	0	-1	4	0	0	0	0	0	0
	5	-1	0	-1	-1	0	5	0	0	0	0	0	0
							$Q =$	0	0	0	0	0	0
							1	0	0	0	0	0	100
							2	0	0	0	0	0	0
							3	0	80	0	0	0	0
							4	0	0	0	0	0	0
							5	0	0	0	0	0	0

References

- Tom Markiewicz & Josh Zheng, Getting started with Artificial Intelligence, Published by O'Reilly Media, 2017
- Stuart J. Russell and Peter Norvig, Artificial Intelligence A Modern Approach
- Richard Szeliski, Computer Vision: Algorithms and Applications, Springer 2010